

بِسْمِ اللّٰهِ الرَّحْمٰنِ الرَّحِيْمِ



An Introduction to Cloud Computing

Session 1

Course Objectives & Description

- The course familiarizes students with cloud computing technologies and practices into a series of well-defined concepts, models, technology mechanisms, and technology architectures, all from an industry-centric and vendor-neutral point of view.
- In this course we will also cover some concept related to Cloud Storage, Security, Load balancing etc. at the same time the case studies will provide a real world, practical perspective on leveraging cloud computing in an organization.

Course Contents

- Fundamentals of Cloud Computing
- Cloud Enabling Technologies
- Cloud Applications
- Popular Cloud Platforms in Industry (AWS, Google AppEngine, Windows Azure)
- Virtualization and Containerization in Cloud Computing
- Web API Services
- Accessing and Manipulating Cloud Data
- Windows Azure SQL Database
- Cloud Load Balancing
- Data Analysis using Azure-ML
- Cloud Security
- Cloud Computing application life cycle management
- An architecture and applications of IoT, Fog/Edge Computing and Cloud Computing

Course Learning Outcomes

| CLO # | CLO DESCRIPTION | BT | PLO |
|-------|---|----|-----|
| CLO 1 | Define the fundamental concepts of cloud computing. | C1 | P1 |
| CLO 2 | Analyze problem requirements to recognize what type of data and processes are involved in a cloud computing solution. | C2 | P2 |
| CLO 3 | Apply and develop business centric models on cloud based IT resources. | C3 | P3 |
| CLO4 | Present the updated technologies evolving in the field of Cloud Computing | A2 | P9 |

Text/Reference books

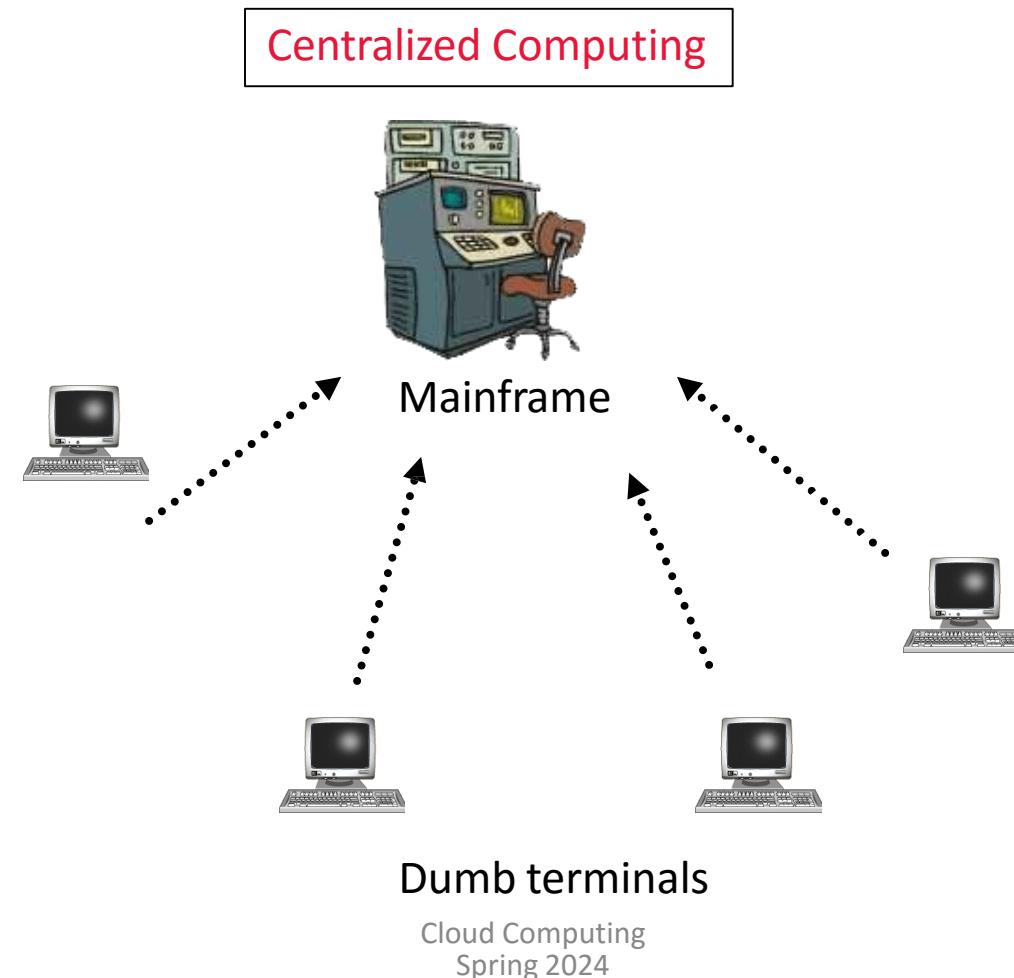
- ***TEXT BOOK***

- Rajkumar Buyya, Christian Vecchiola “Mastering Cloud Computing”, McGraw Hill
- THOMAS ERL, “CLOUD COMPUTING CONCEPTS, TECHNOLOGY AND ARCHITECTURE”, 1ST ED., 2014, PEARSON.

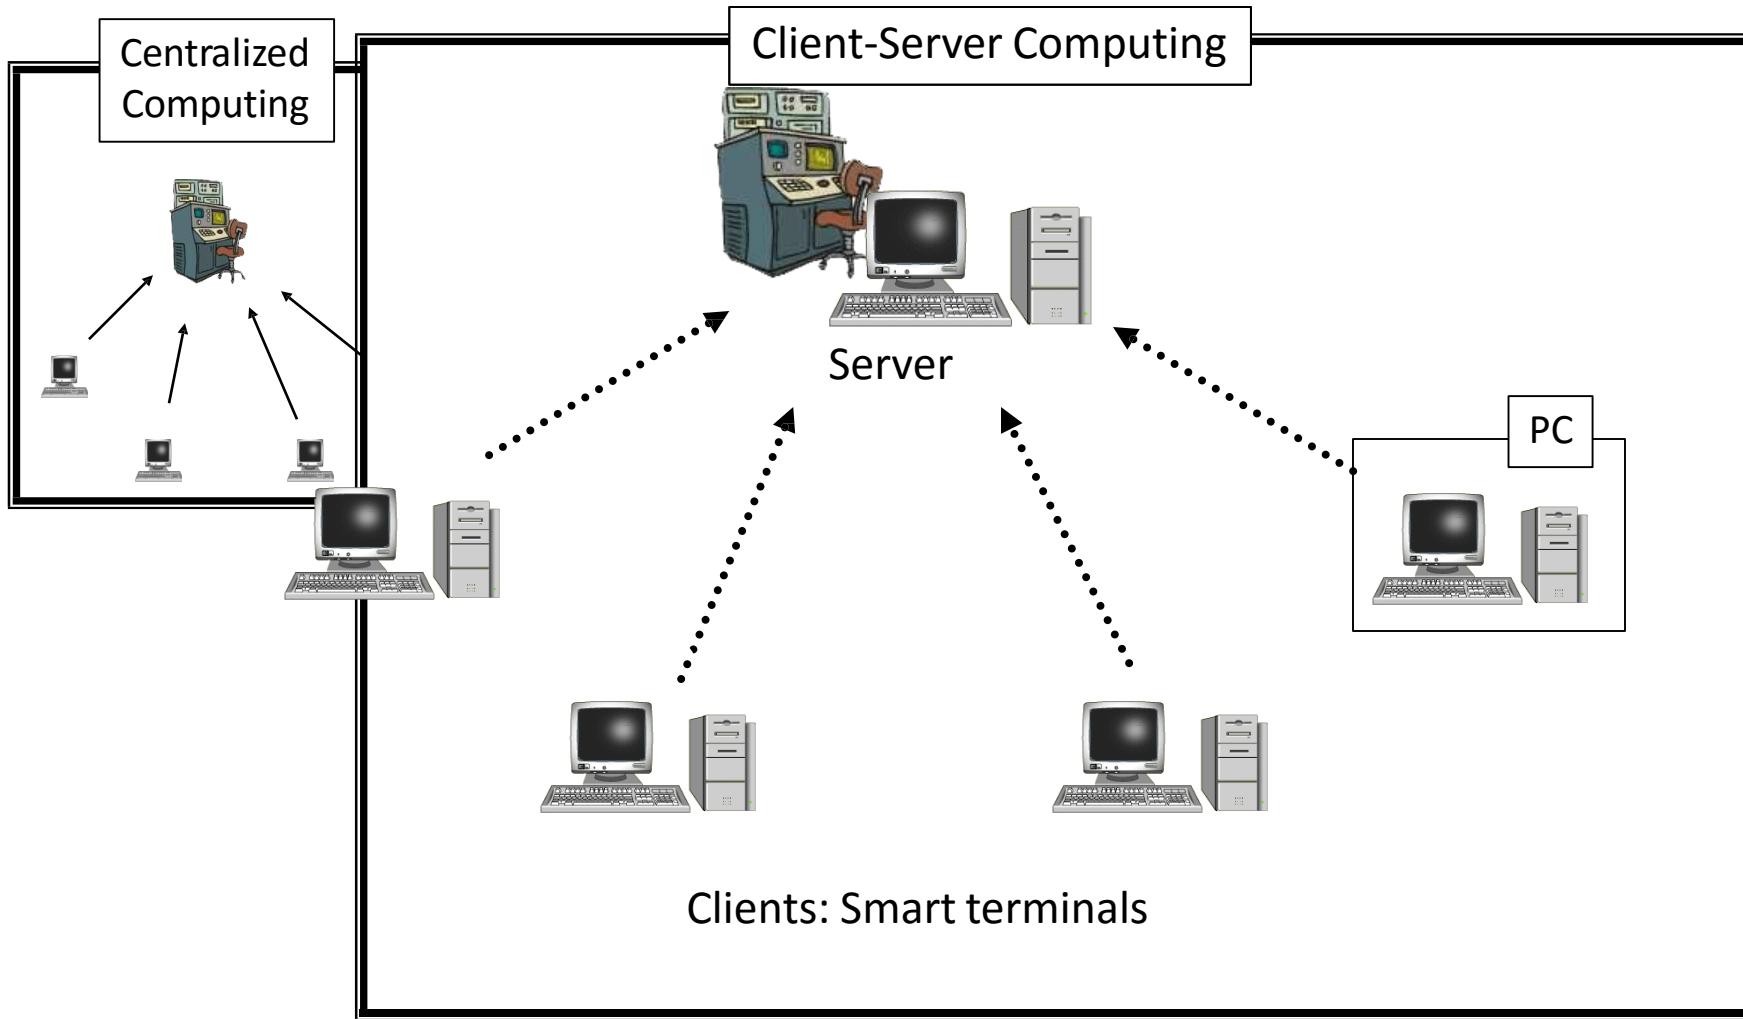
- ***REFERENCE BOOKS***

- RONALD L. KRUTZ AND RUSSELL DEAN VINES, “CLOUD SECURITY-A COMPREHENSIVE GUIDE TO SECURE CLOUD COMPUTING”, 2013, WILEY.
- BARRIE SOSINKY, “CLOUD COMPUTING”, 1ST ED., 2014, WILEY.
- “Cloud Computing: Principles and Paradigms”, John Wiley & Sons

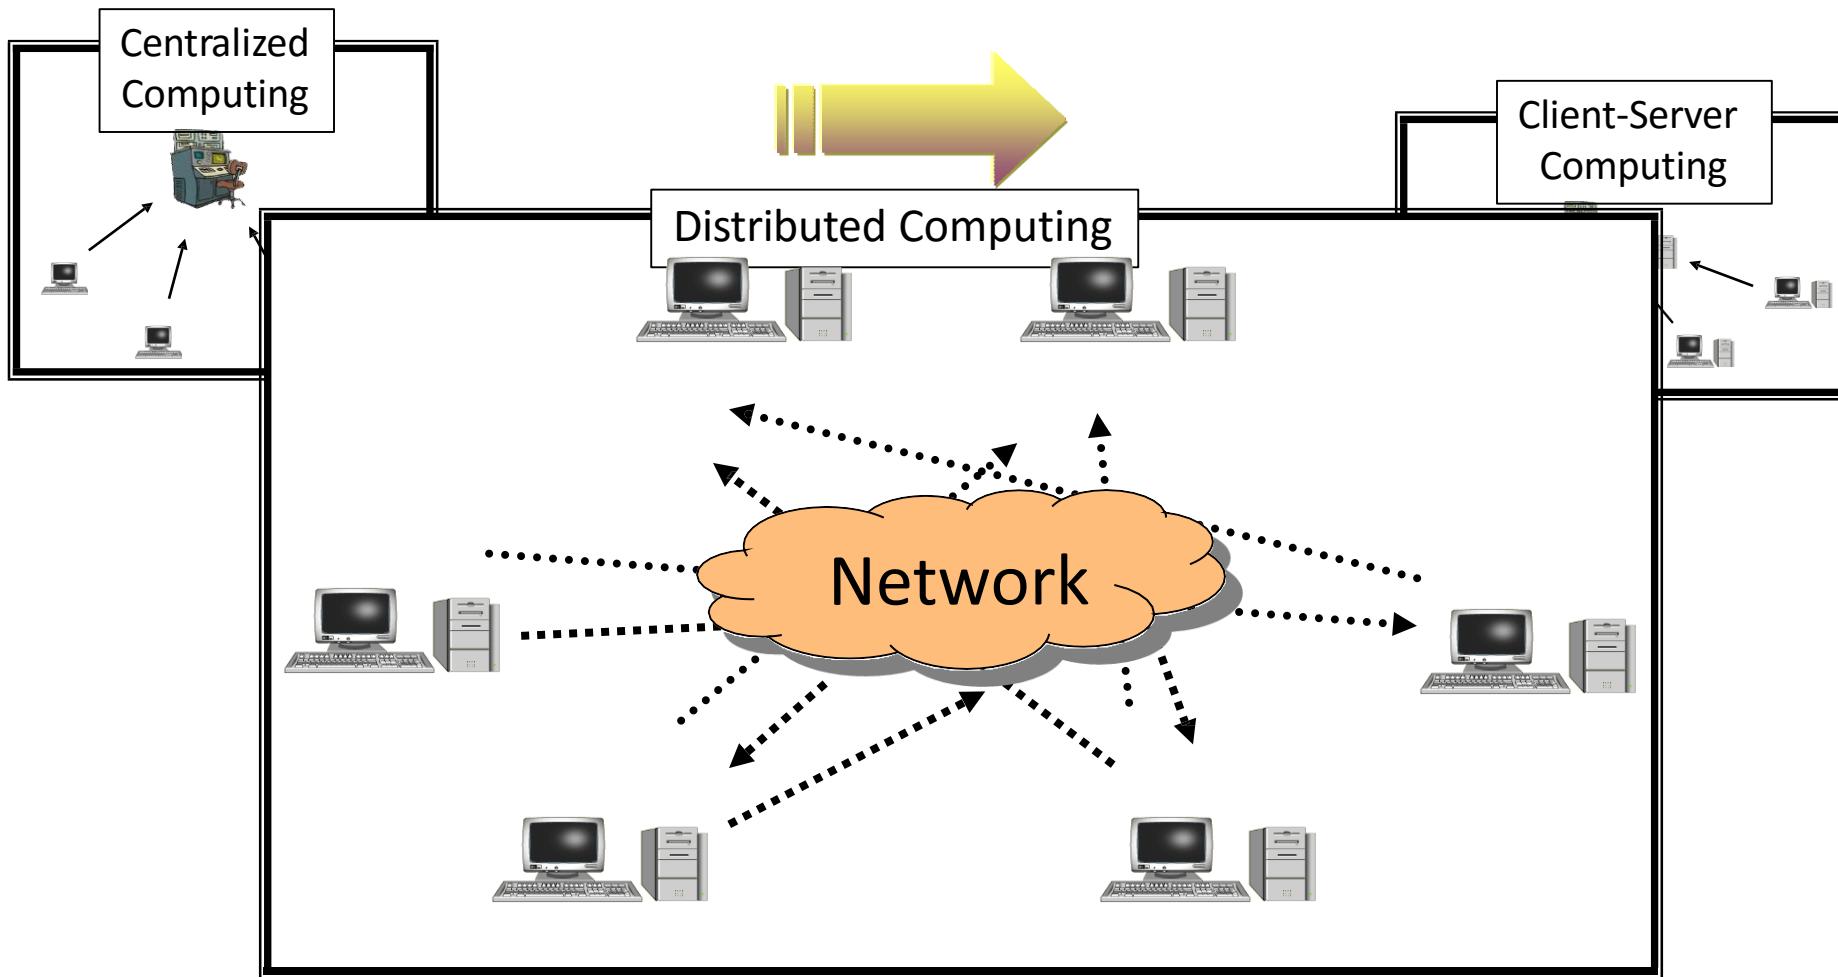
Transformation in Computing



Transformation in Computing

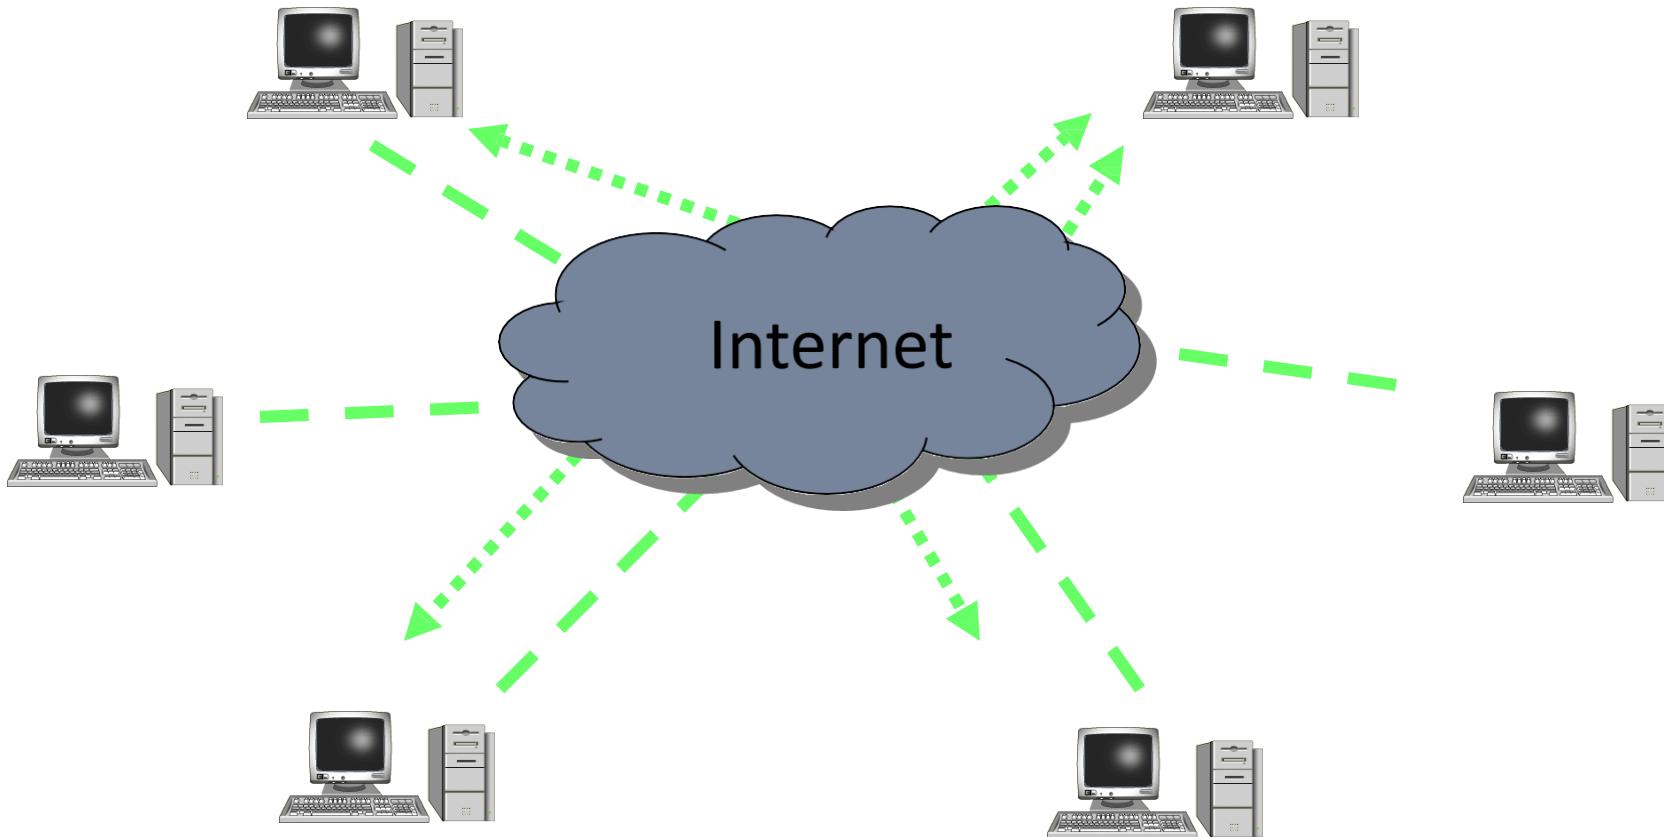


Transformation in Computing

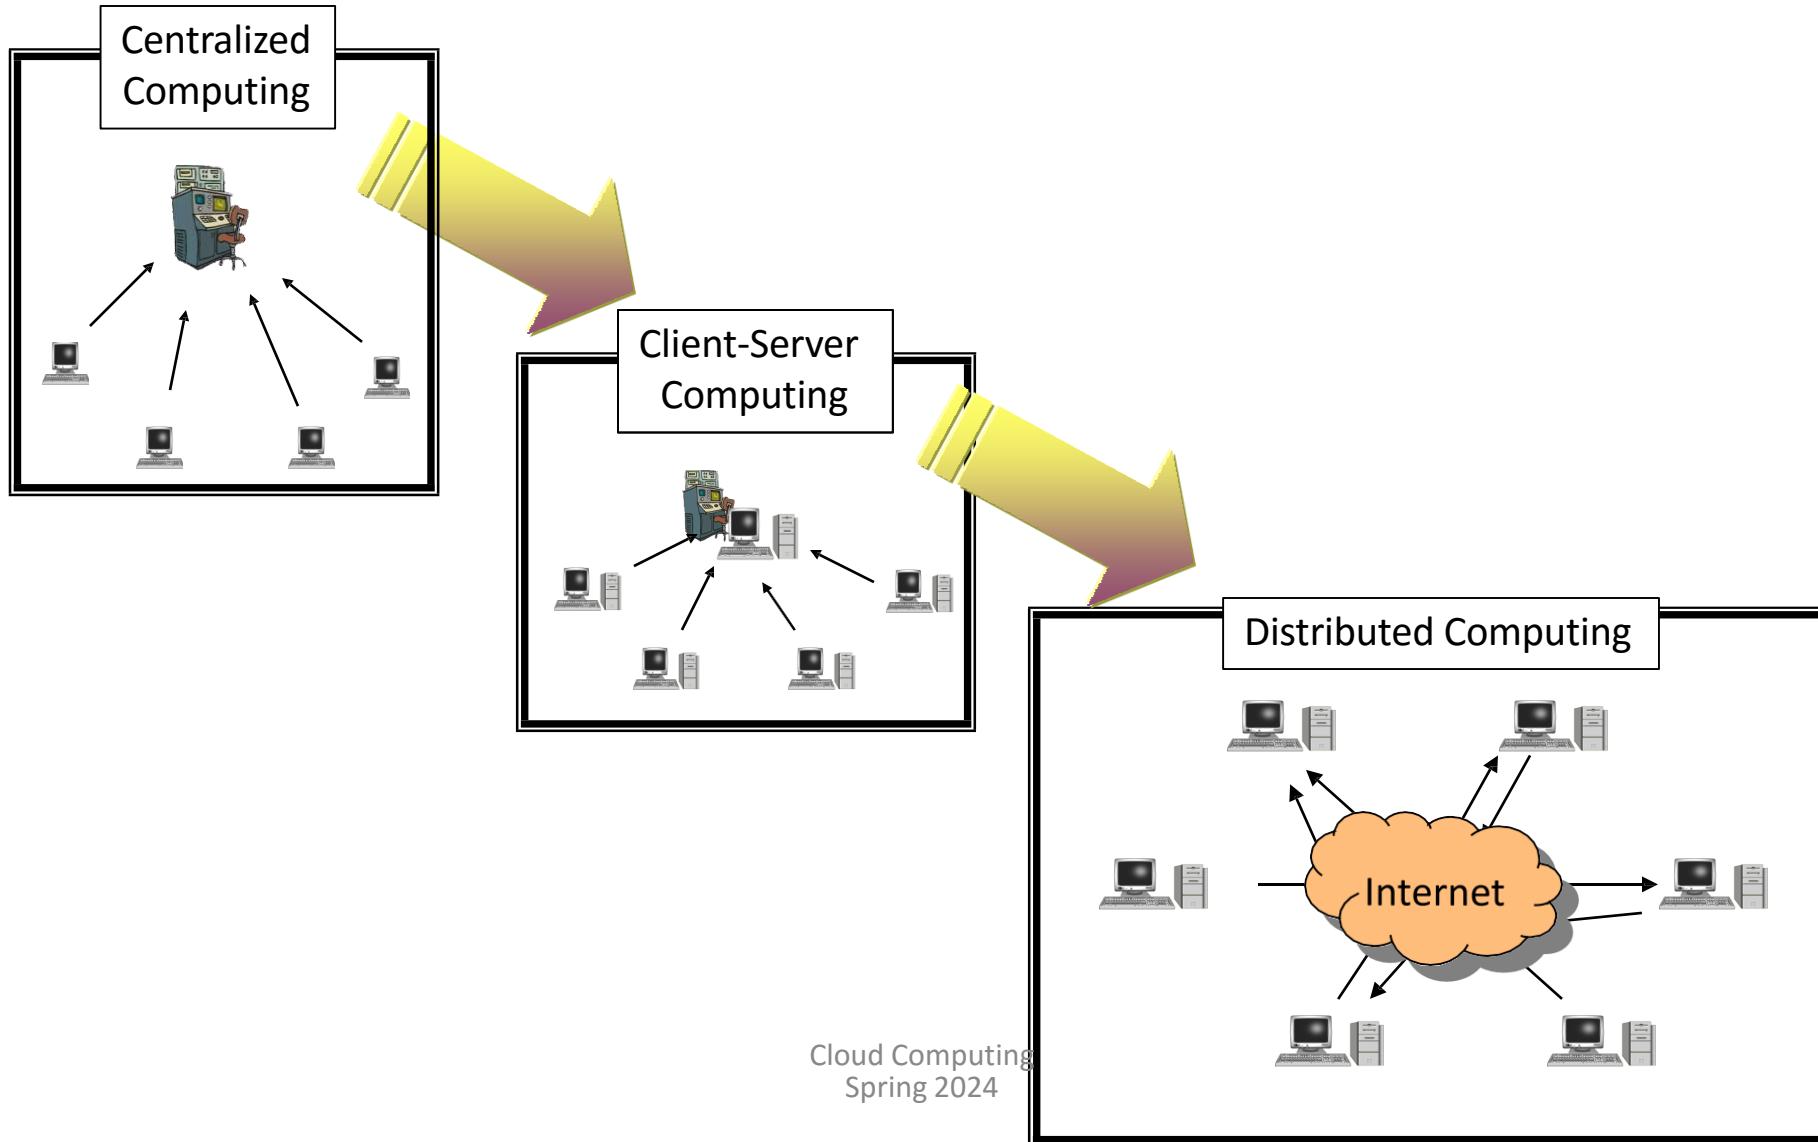


In local distributed computing, computers are present generally in a LAN

Distributed Computing

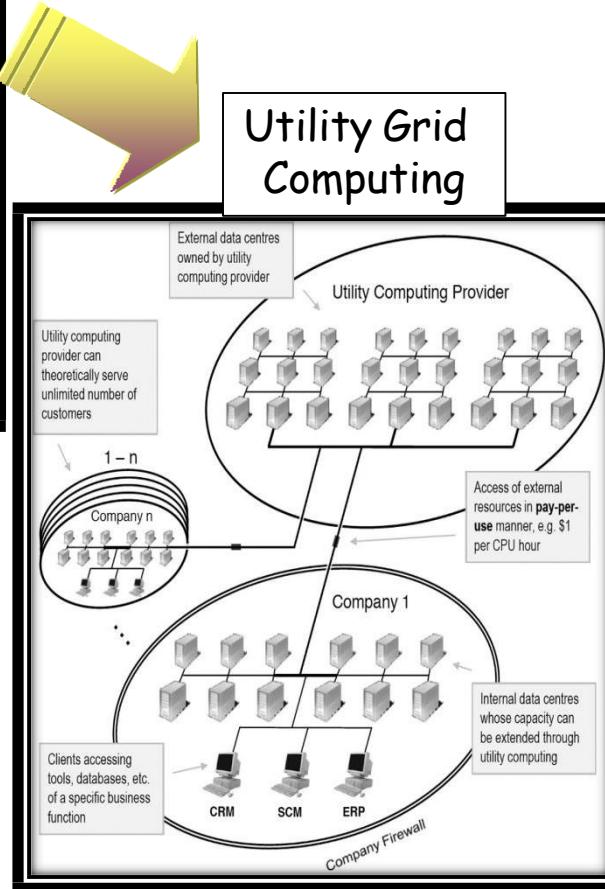
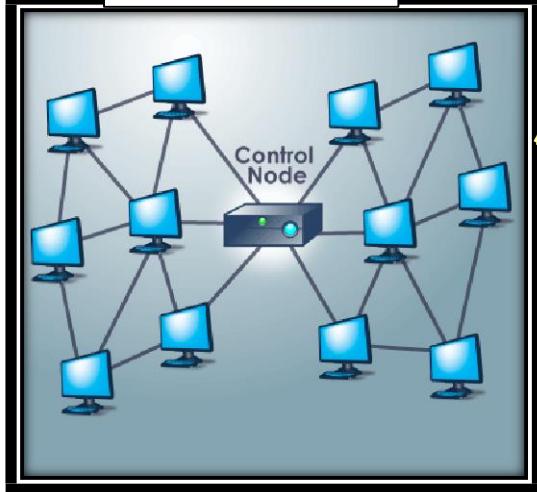


Transformation in Computing

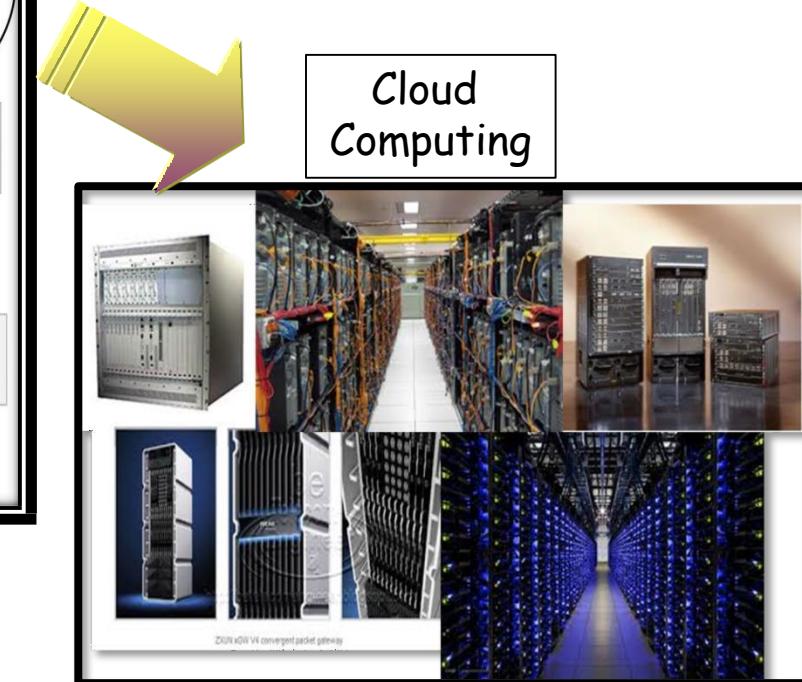


Cont....

Cluster/Grid Computing



Cloud Computing
Spring 2024



What is a Cloud?

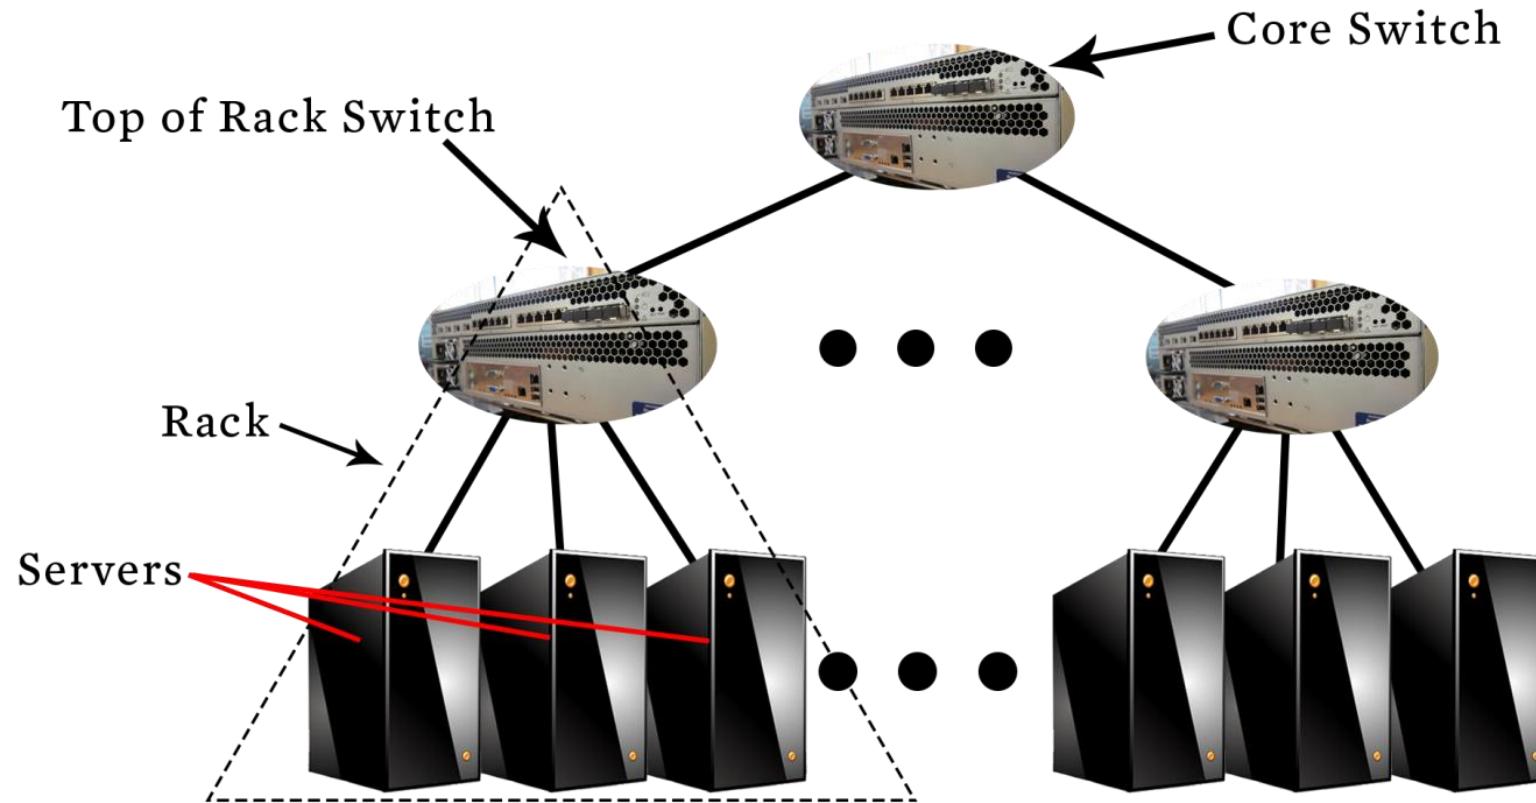
- It's a cluster!
- It's a supercomputer!
- It's a datastore!
- It's superman!
- None of the above
- All of the above
- Cloud = **Lots of storage + compute cycles nearby**



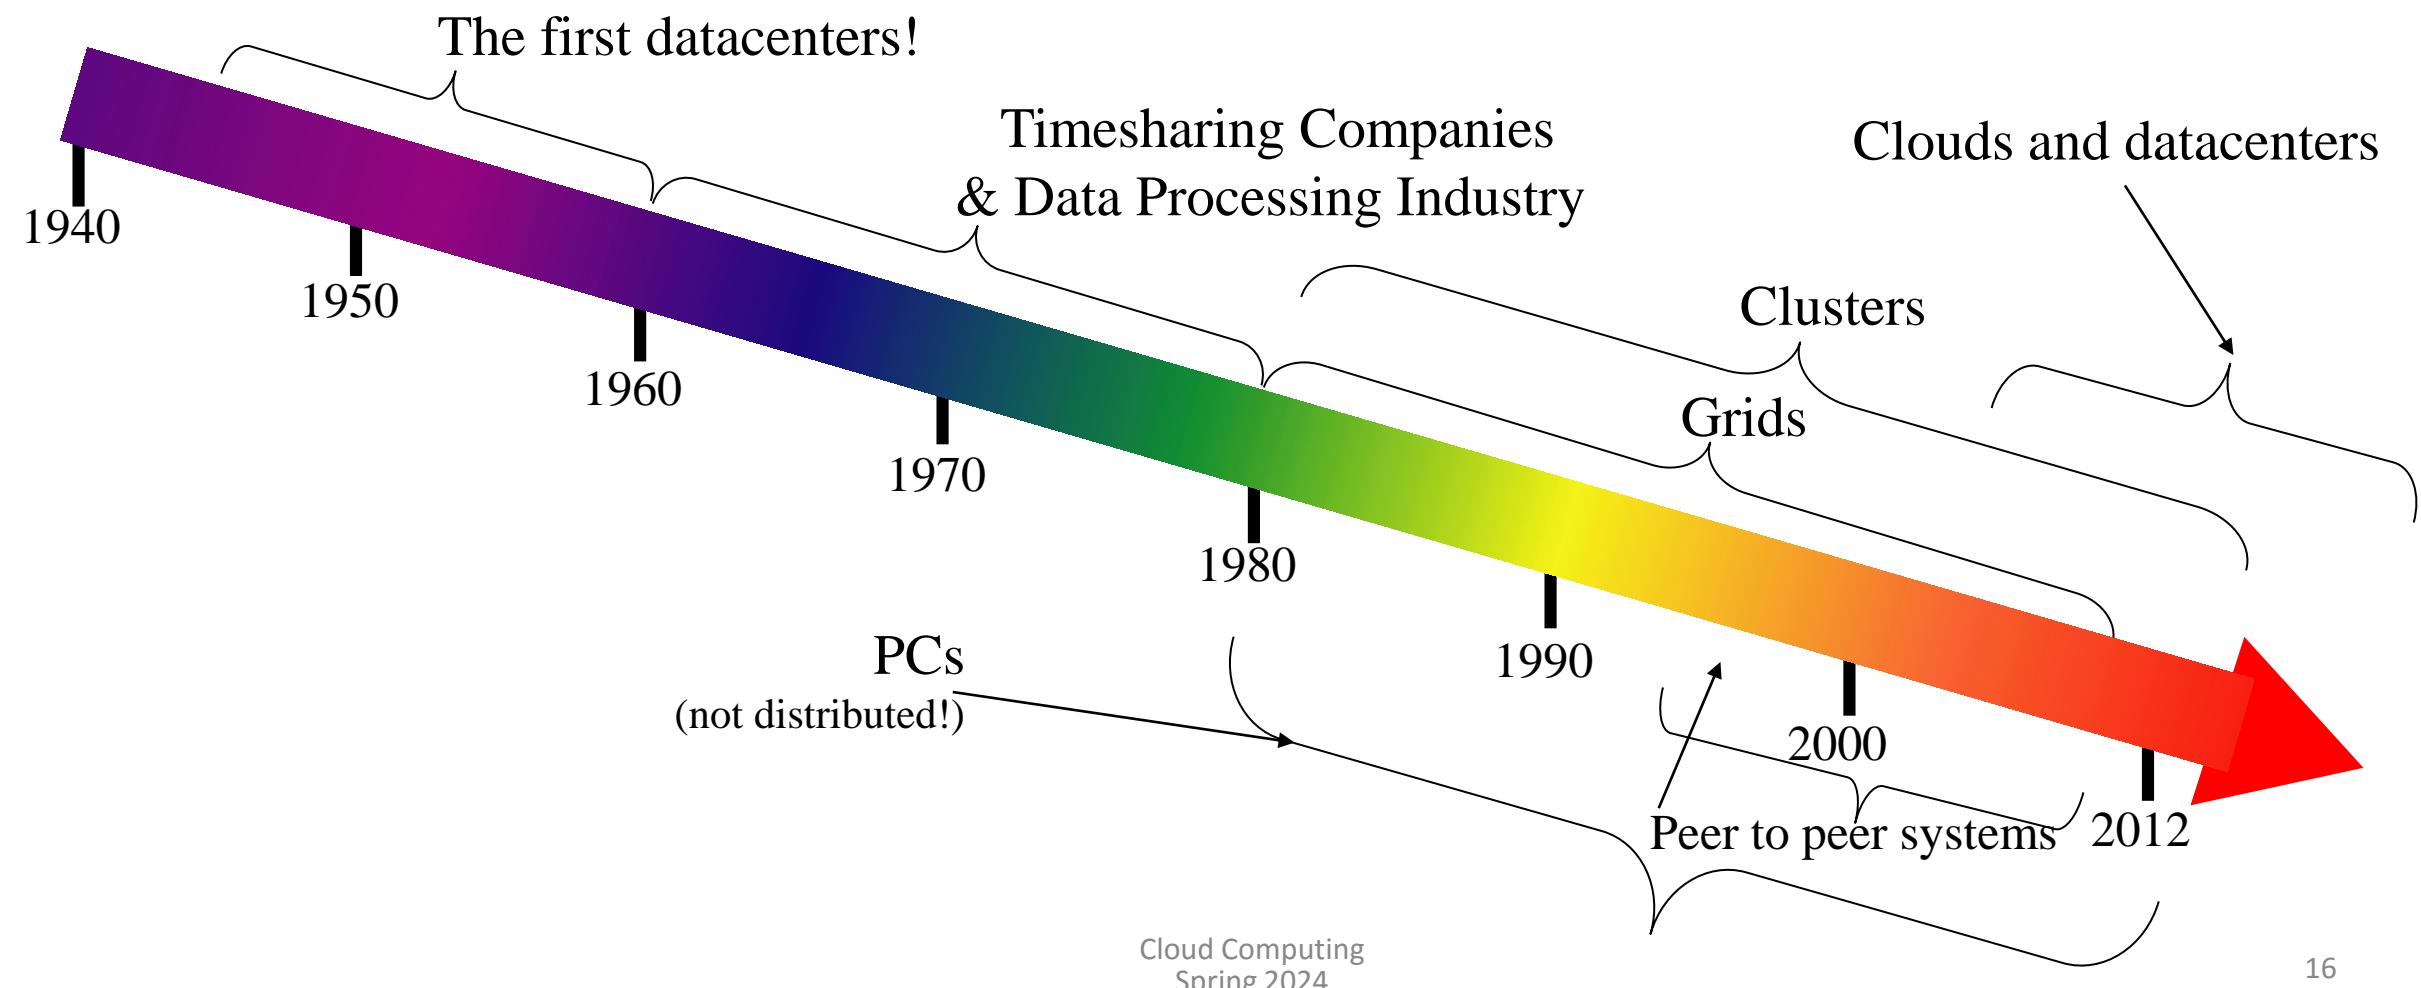
What is a Cloud?

- A single-site cloud (aka “Datacenter”) consists of
 - Compute nodes (grouped into racks) (2)
 - Switches, connecting the racks
 - A network topology, e.g., hierarchical
 - Storage (backend) nodes connected to the network (3)
 - Front-end for submitting jobs and receiving client requests (1)
 - (1-3: Often called “three-tier architecture”)
 - Software Services
- A geographically distributed cloud consists of
 - Multiple such sites
 - Each site perhaps with a different structure and services

A Sample Cloud Topology



“A Cloudy History of Time”



Servers



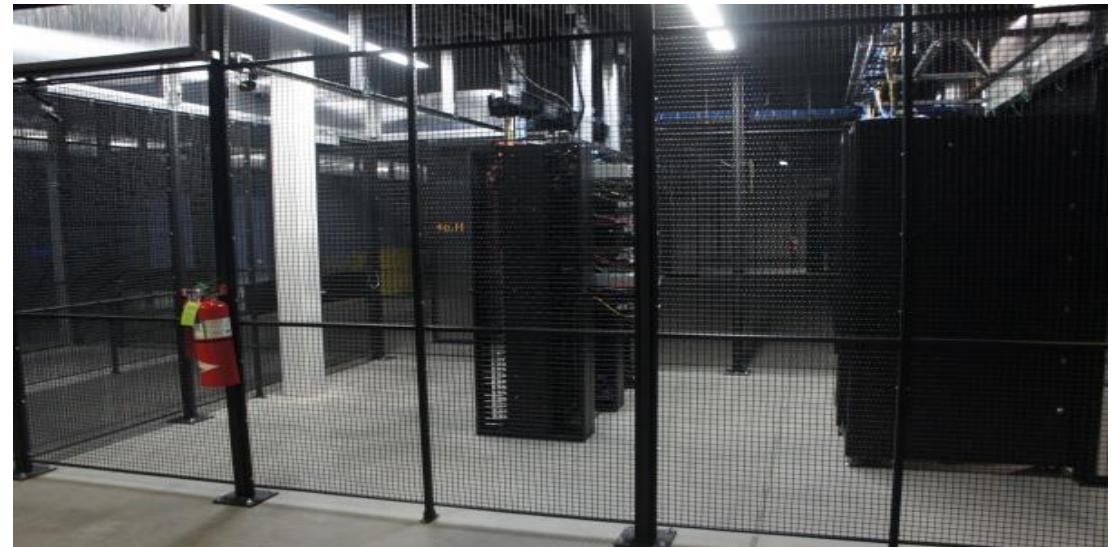
Front



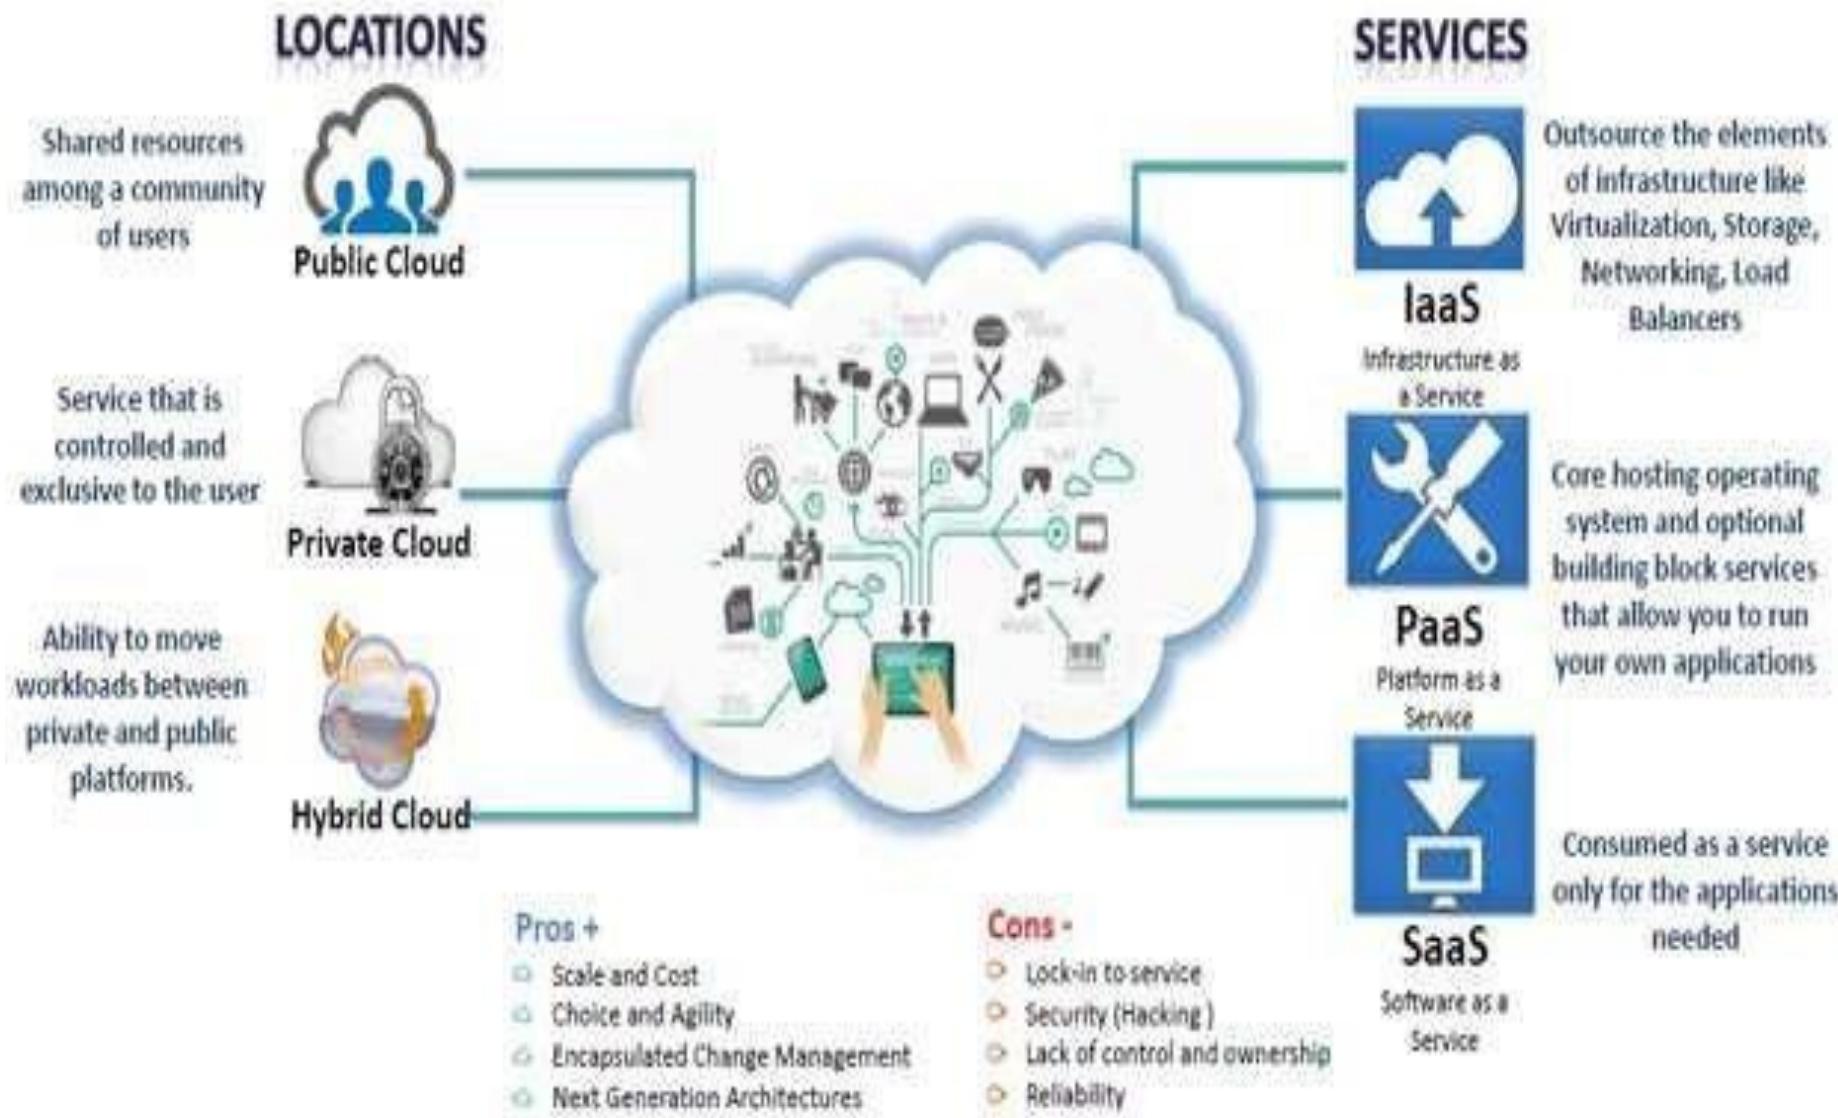
In



Back



What is Cloud Computing?



The Next Revolution in IT

The Big Switch in IT

- Classical Computing
 - Buy & Own
 - Hardware, System Software, Applications often to meet peak needs.
 - Install, Configure, Test, Verify
 - Manage
 - ..
 - Finally, use it
 - – \$\$\$\$\$....\$(High CapEx)



- Cloud Computing
 - Subscribe
 - Use
- based on QoS



Every 18 months?

Conventional Computing

vs.

Cloud Computing

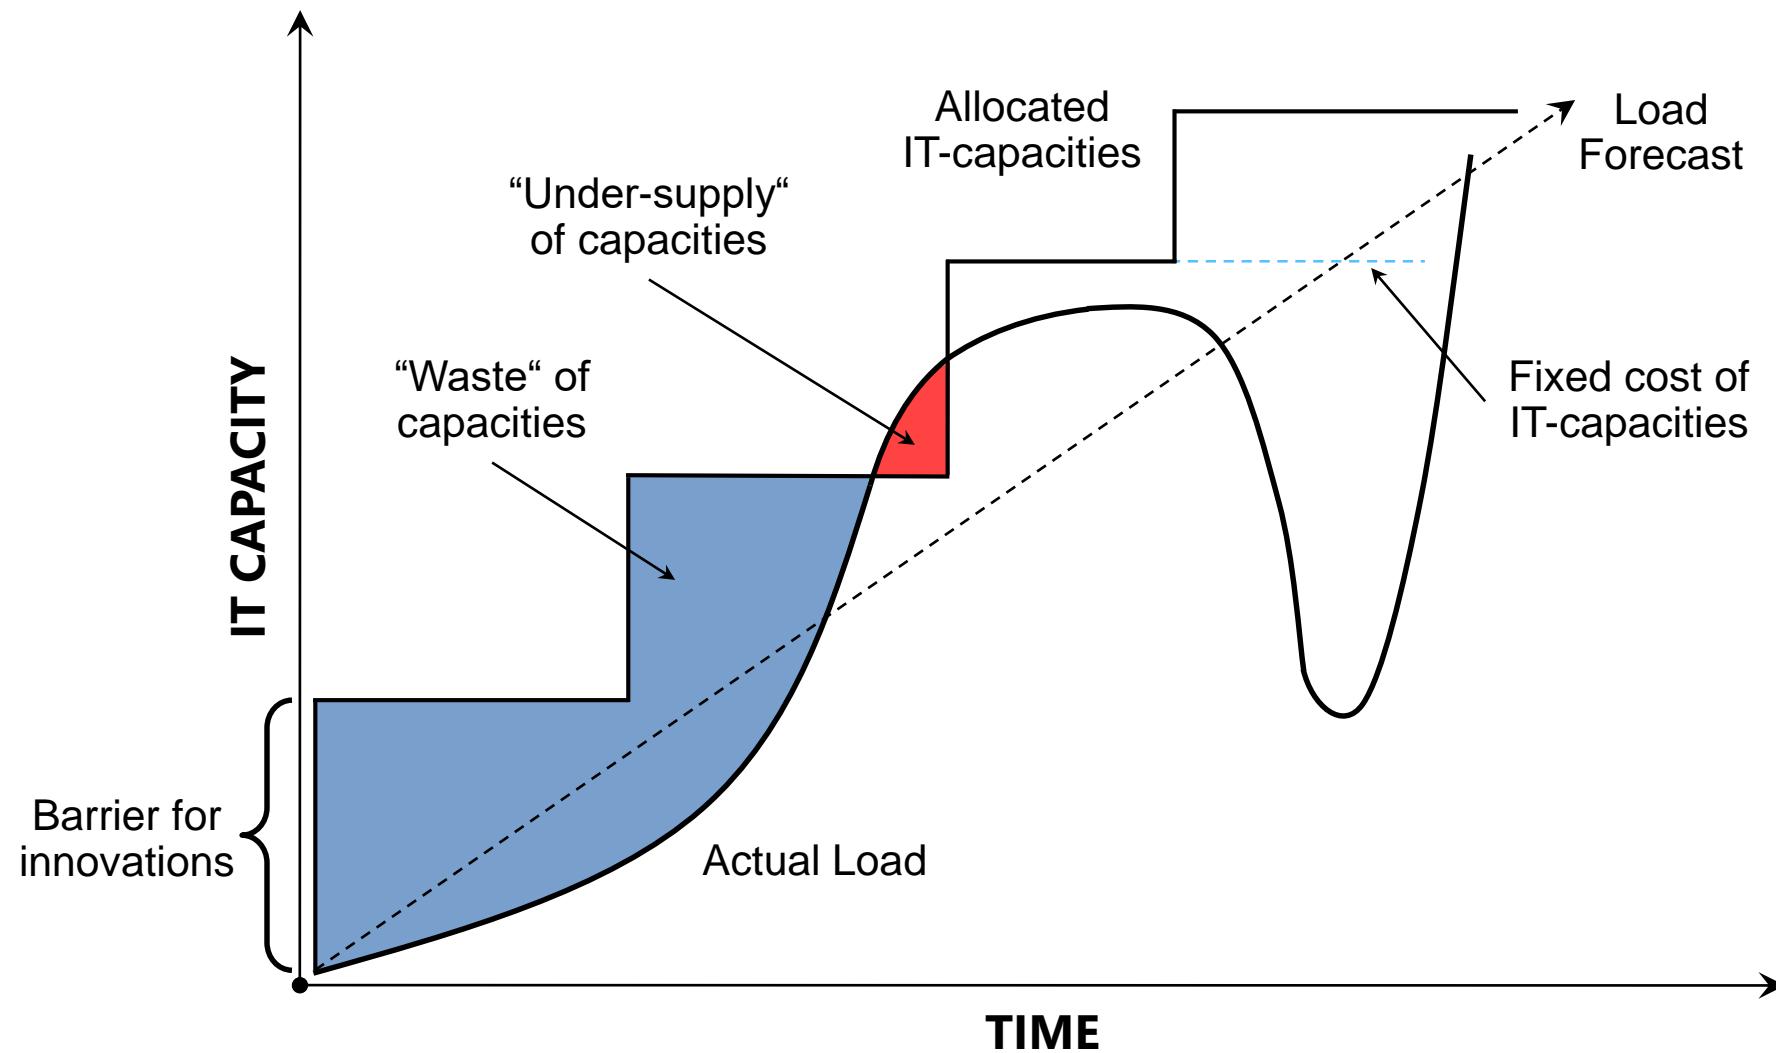
Conventional

- Manually Provisioned
- Dedicated Hardware
- Fixed Capacity
- Pay for Capacity
- Capital & Operational Expenses
- Managed via System administrators

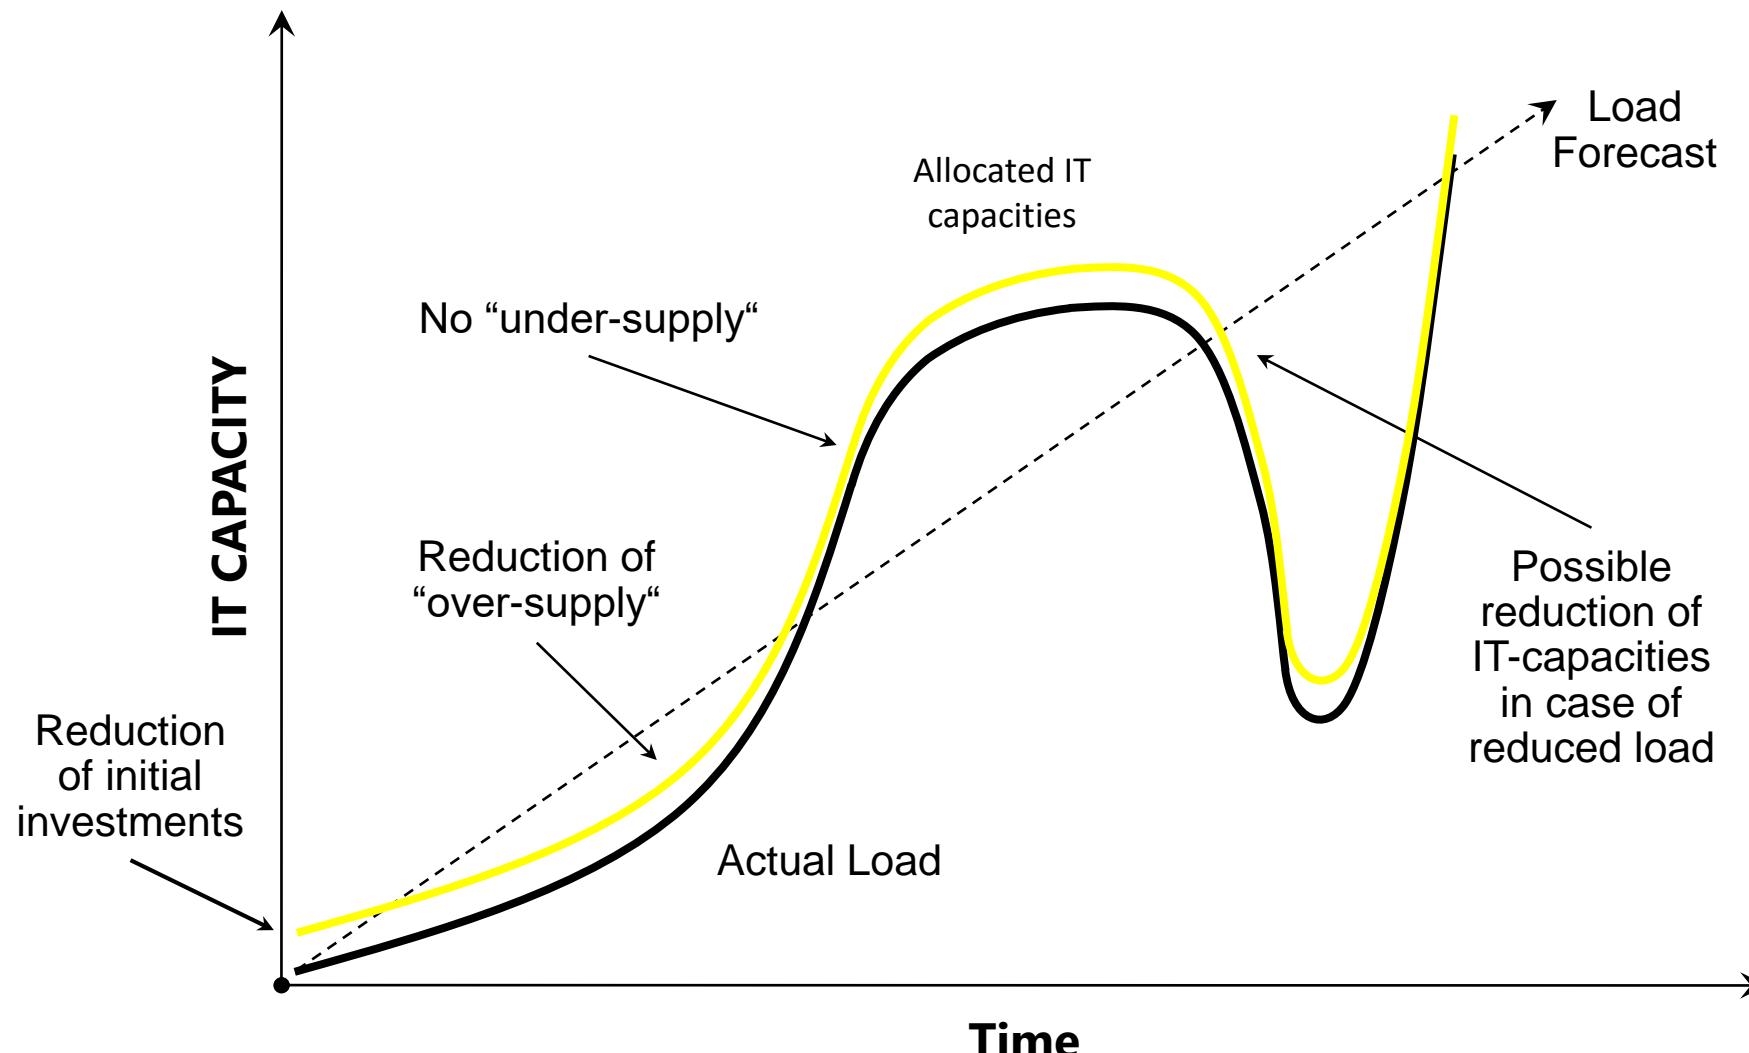
Cloud

- Self-provisioned
- Shared Hardware
- Elastic Capacity
- Pay for Use
- Operational Expenses
- Managed via APIs

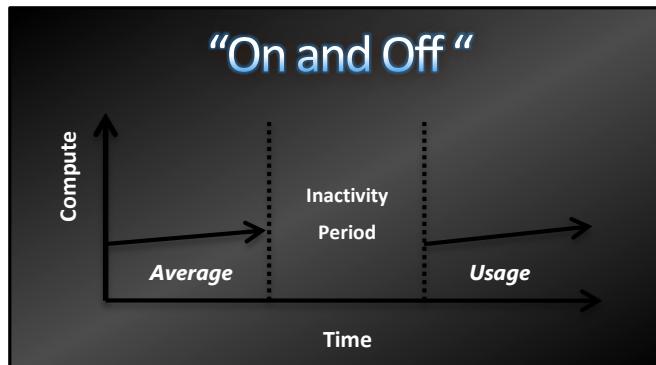
A non-Cloud view



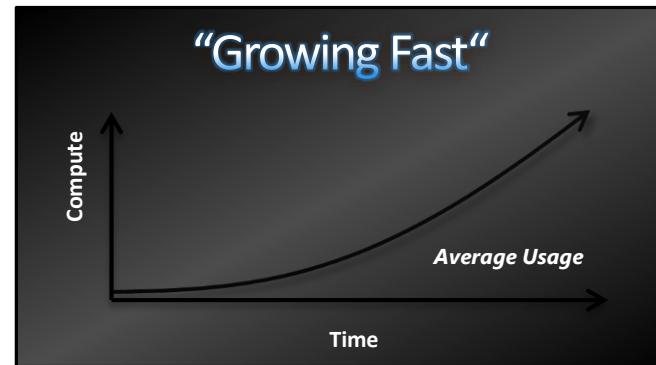
A Cloud View



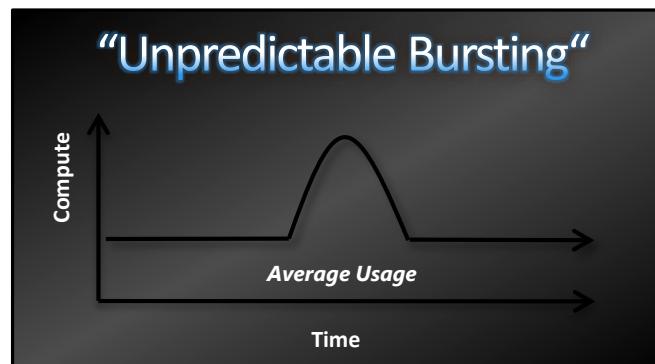
Load Pattern



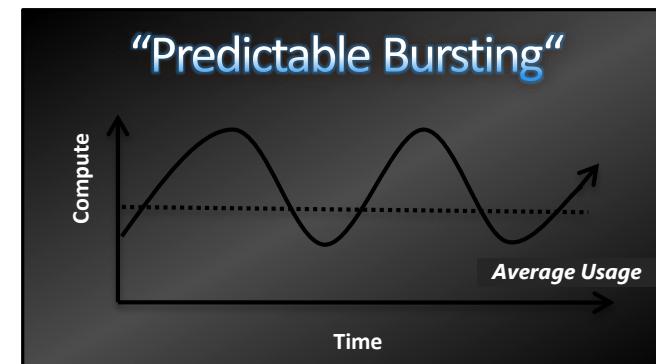
- On & off workloads (e.g. batch job)
- Over provisioned capacity is wasted
- Time to market can be cumbersome



- Successful services need to grow/scale
- Keeping up w/ growth is big IT challenge
- Cannot provision hardware fast enough

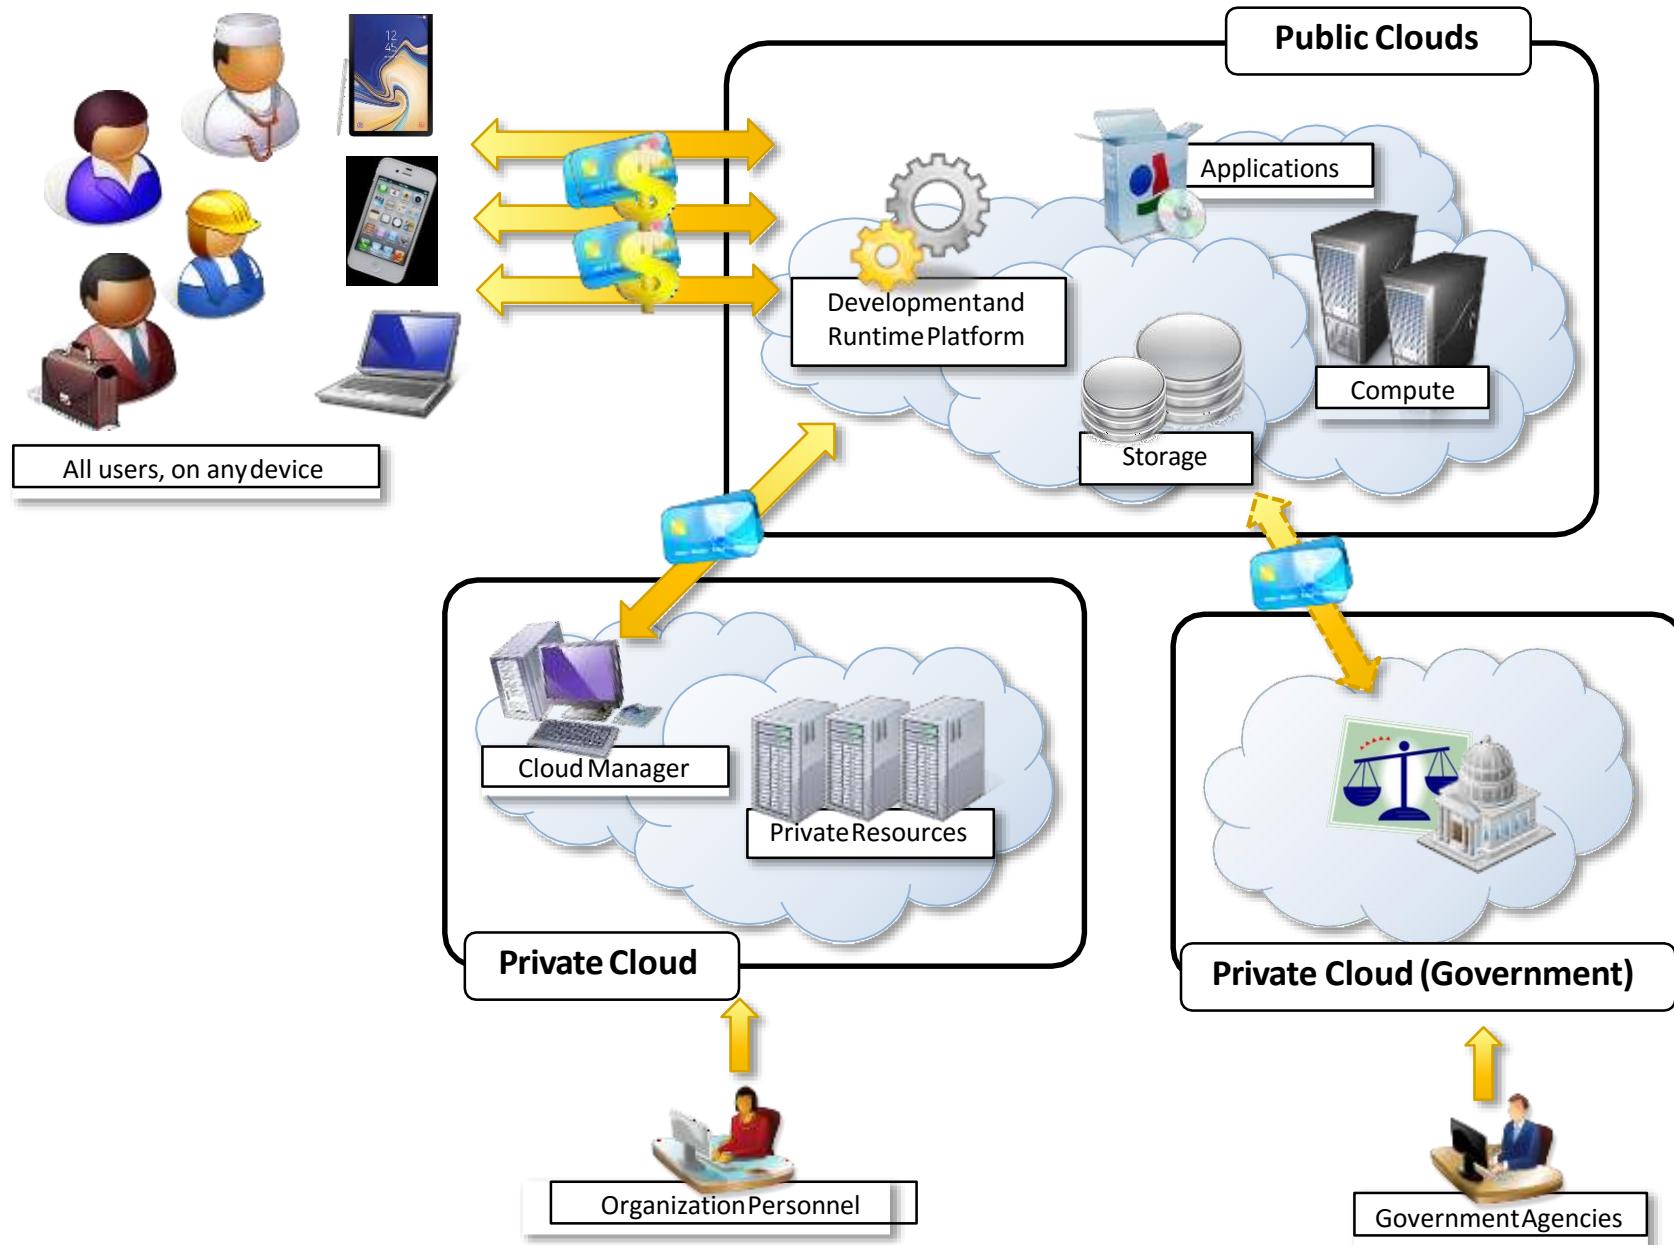


- Unexpected/unplanned peak in demand
- Sudden spike impacts performance
- Can't over provision for extreme cases

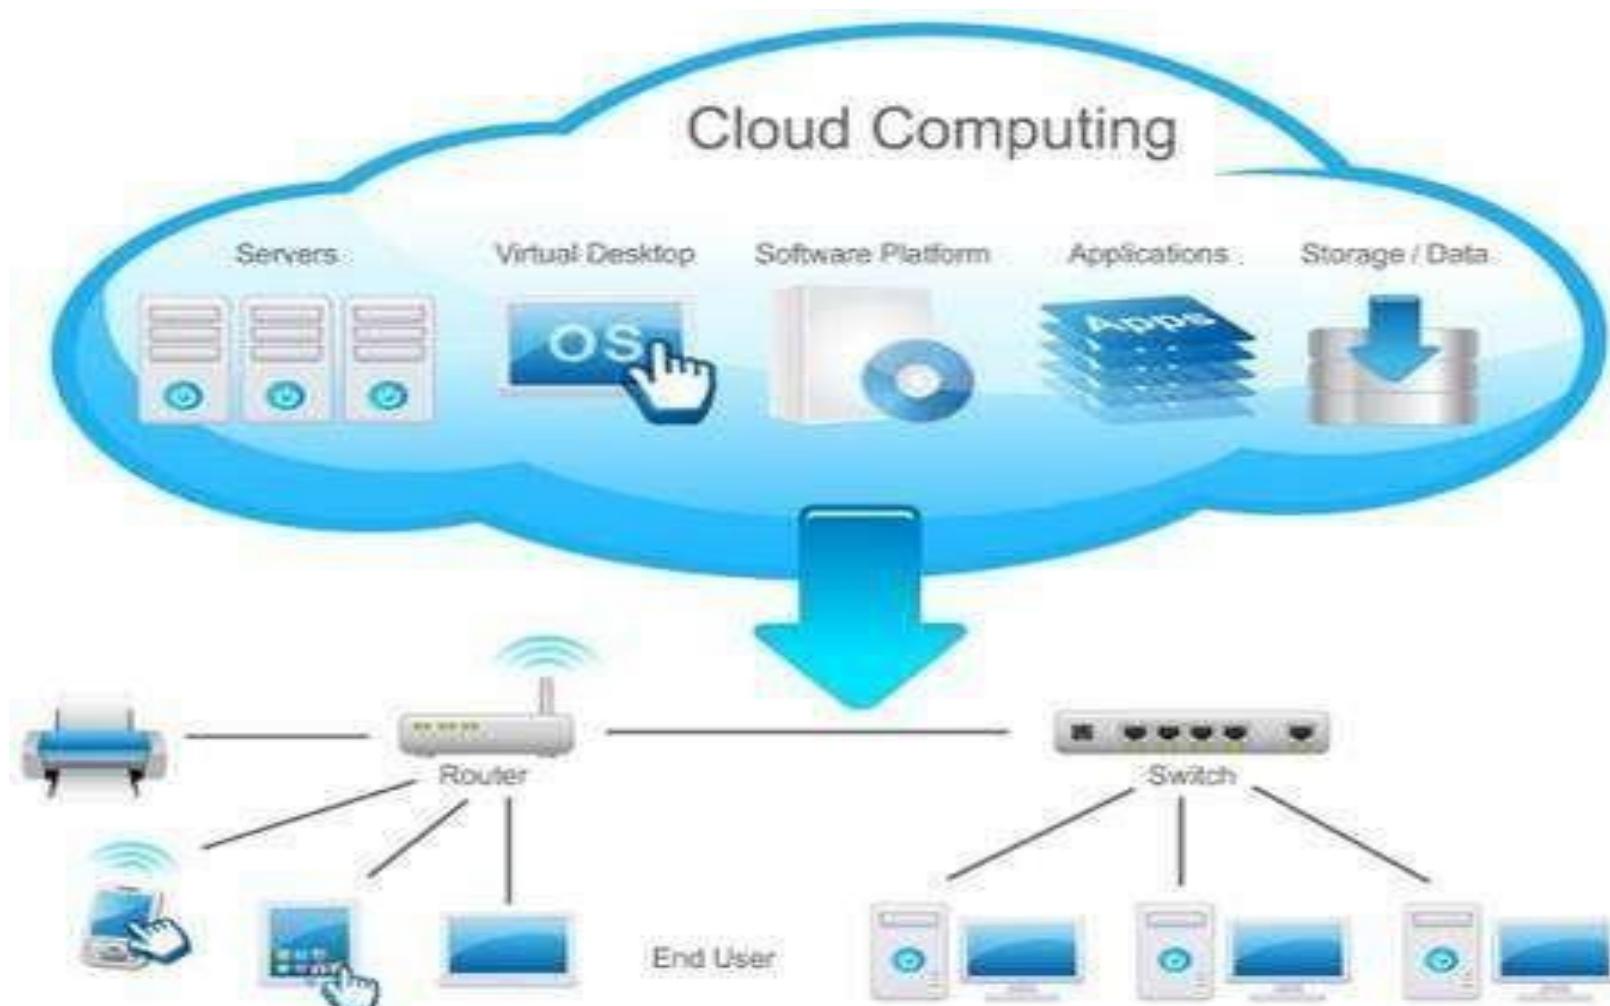


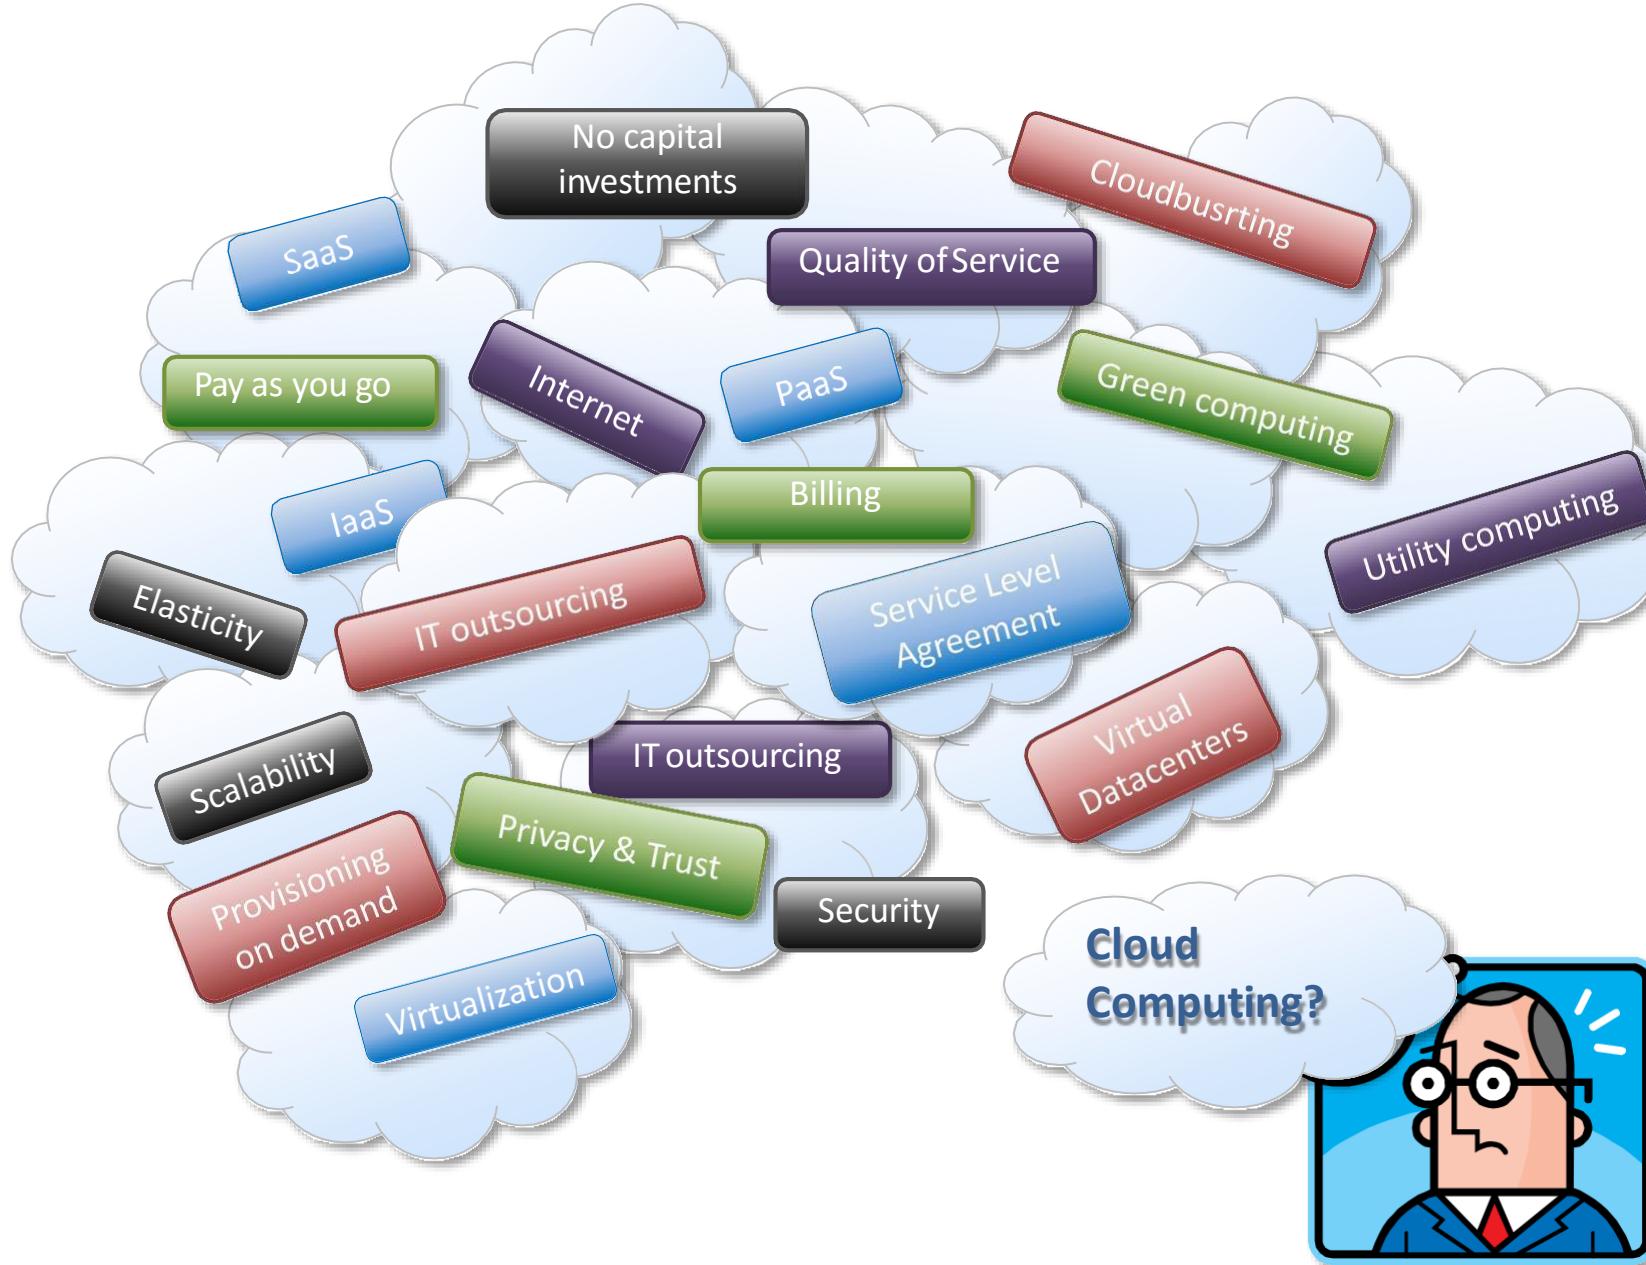
- Services with micro seasonality trends
- Peaks due to periodic increased demand
- IT complexity and wasted capacity





Defining Cloud





Defining Cloud



A Closer look

Cloud computing Helping:

- Enterprises
- Governments
- Public Institutes
- Private Institutes
- Research Organization

Examples

- Large enterprise can offload some of their activities to Cloud based system.



Example

- Small Enterprises and Start-ups can afford to translate into business results their ideas more quickly without excessive upfront cost



Example

- System Developers can concentrate on business logic rather than dealing with the complexity of infrastructure management and scalability

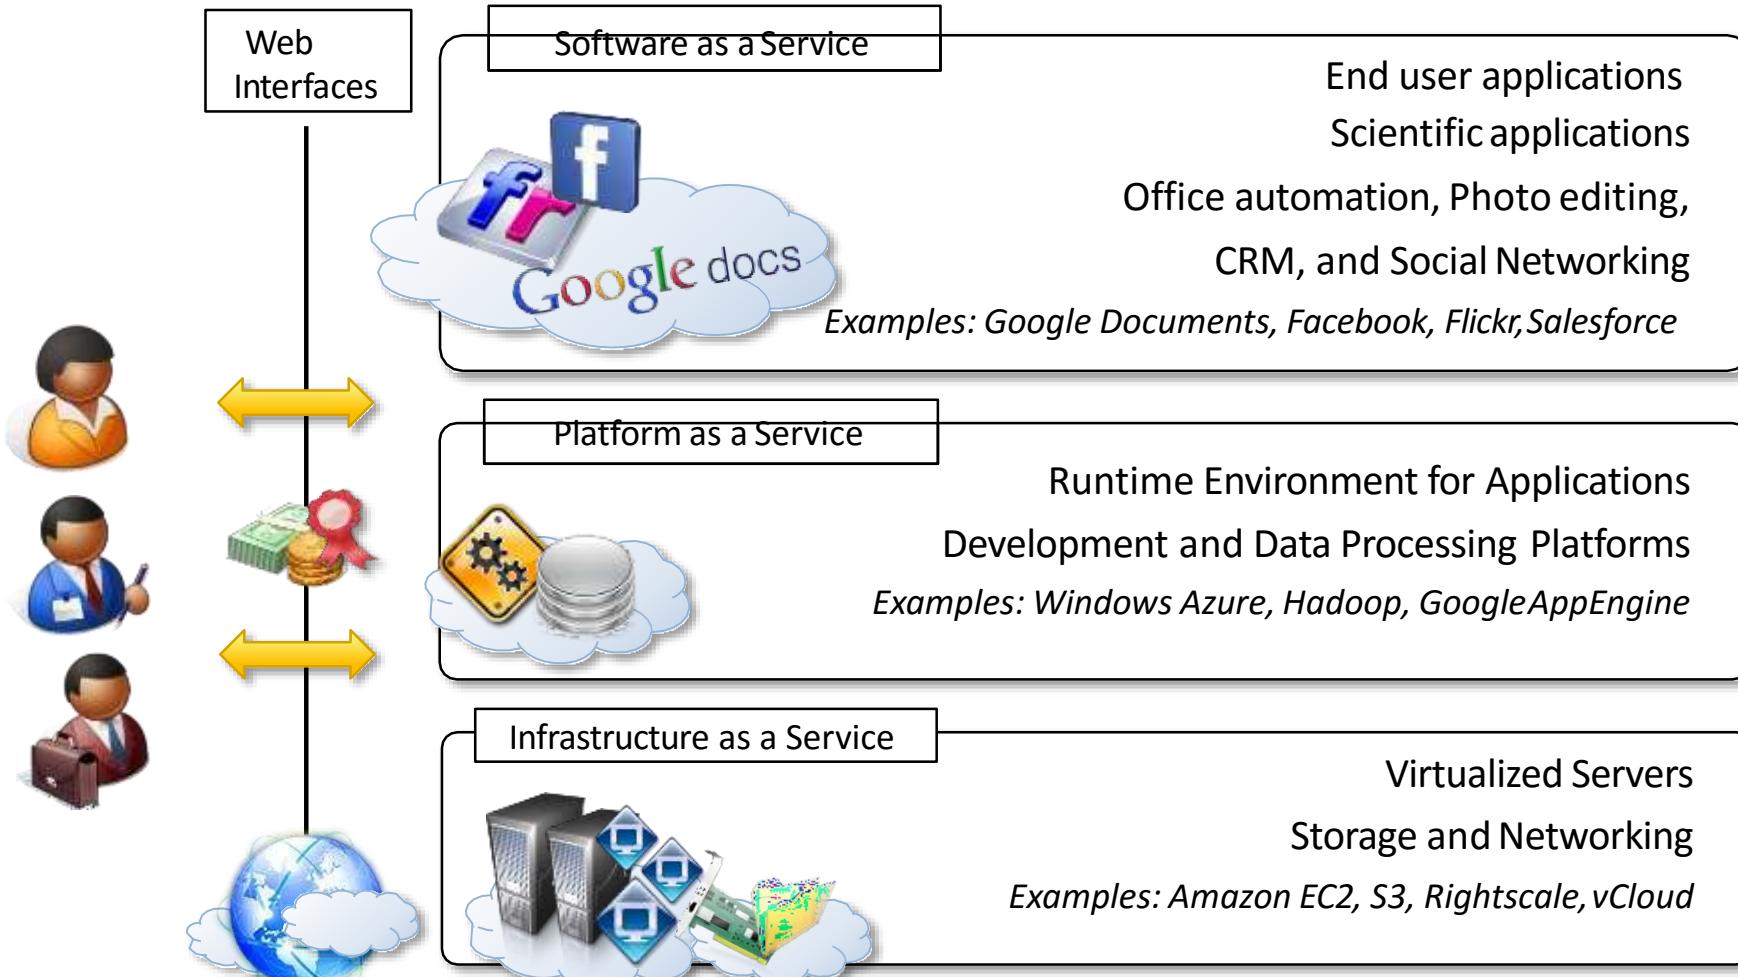


Example

- End users can have their documents accessible from everywhere and any device



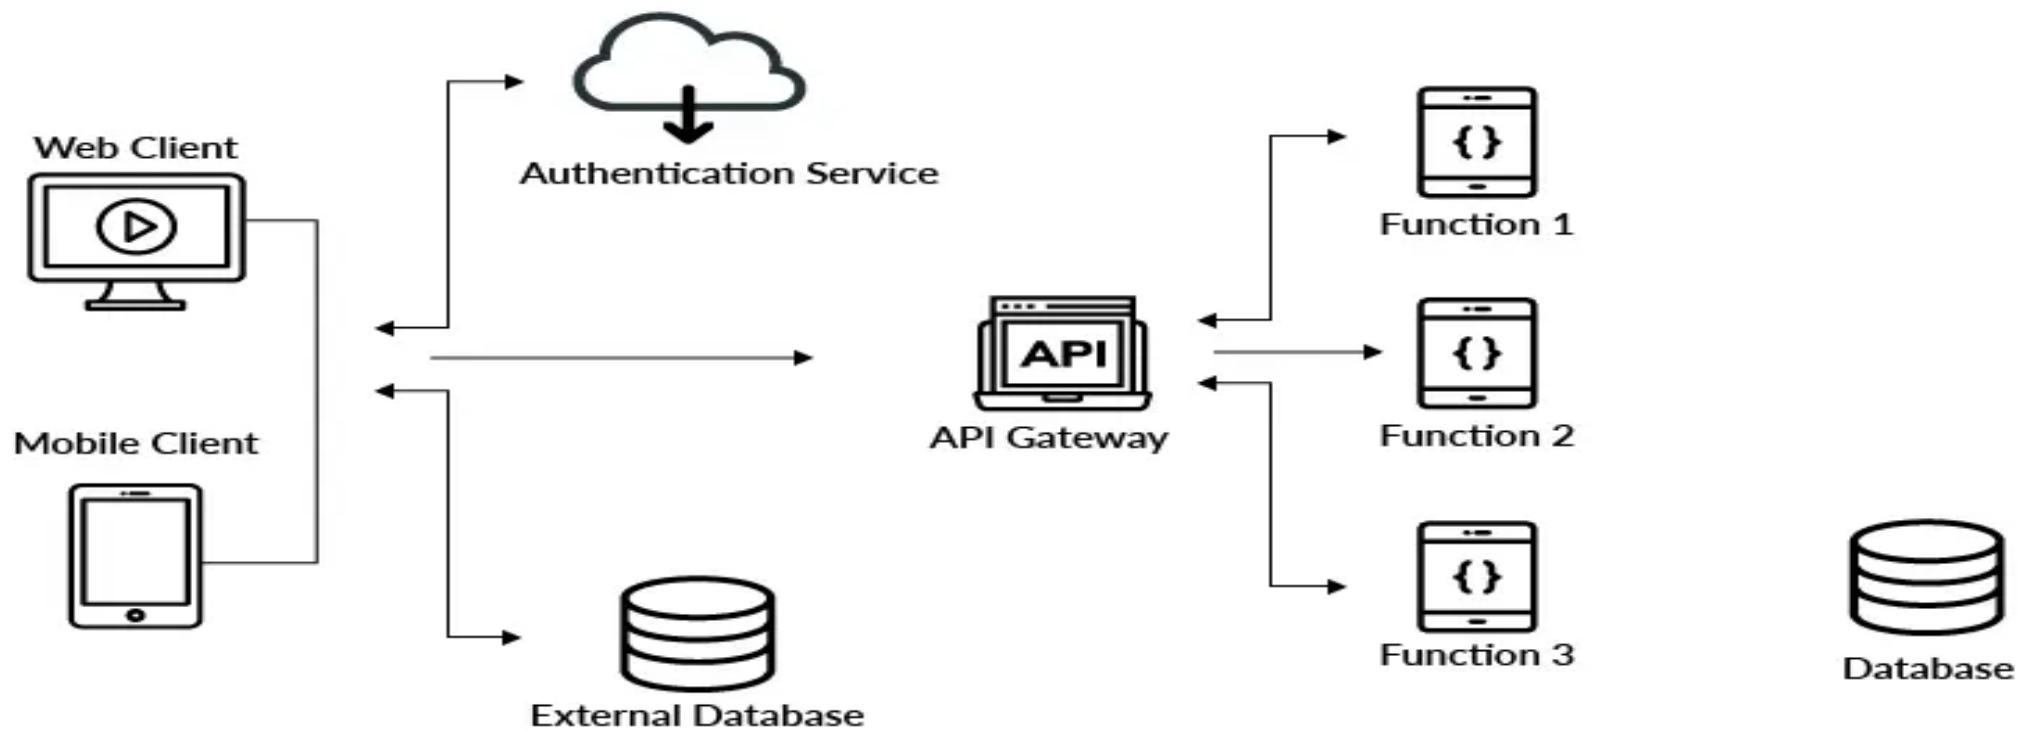
Cloud Computing Reference Model



Serverless Computing

- Serverless computing is a cloud application development and execution model that lets developers build and run code without managing servers, and without paying for idle cloud infrastructure.
- Serverless lets developers put all their focus into writing the best front-end application code and business logic they can.
- All developers need to do is write their application code and deploy it to containers managed by a cloud service provider.
- The cloud provider handles the rest, provisioning the cloud infrastructure required to run the code and scaling the infrastructure up and down on demand as needed.
- The cloud provider is also responsible for all routine infrastructure management and maintenance such as operating system updates and patches, security management, capacity planning, system monitoring and more.
- With serverless, developers never pay for idle capacity. The cloud provider spins up and provisions the required computing resources on demand when the code executes, and spins them back down again—called ‘scaling to zero’—when execution stops.
- The billing starts when execution starts, and ends when execution stops; typically, pricing is based on execution time and resources required.

Working Of Serverless Architecture



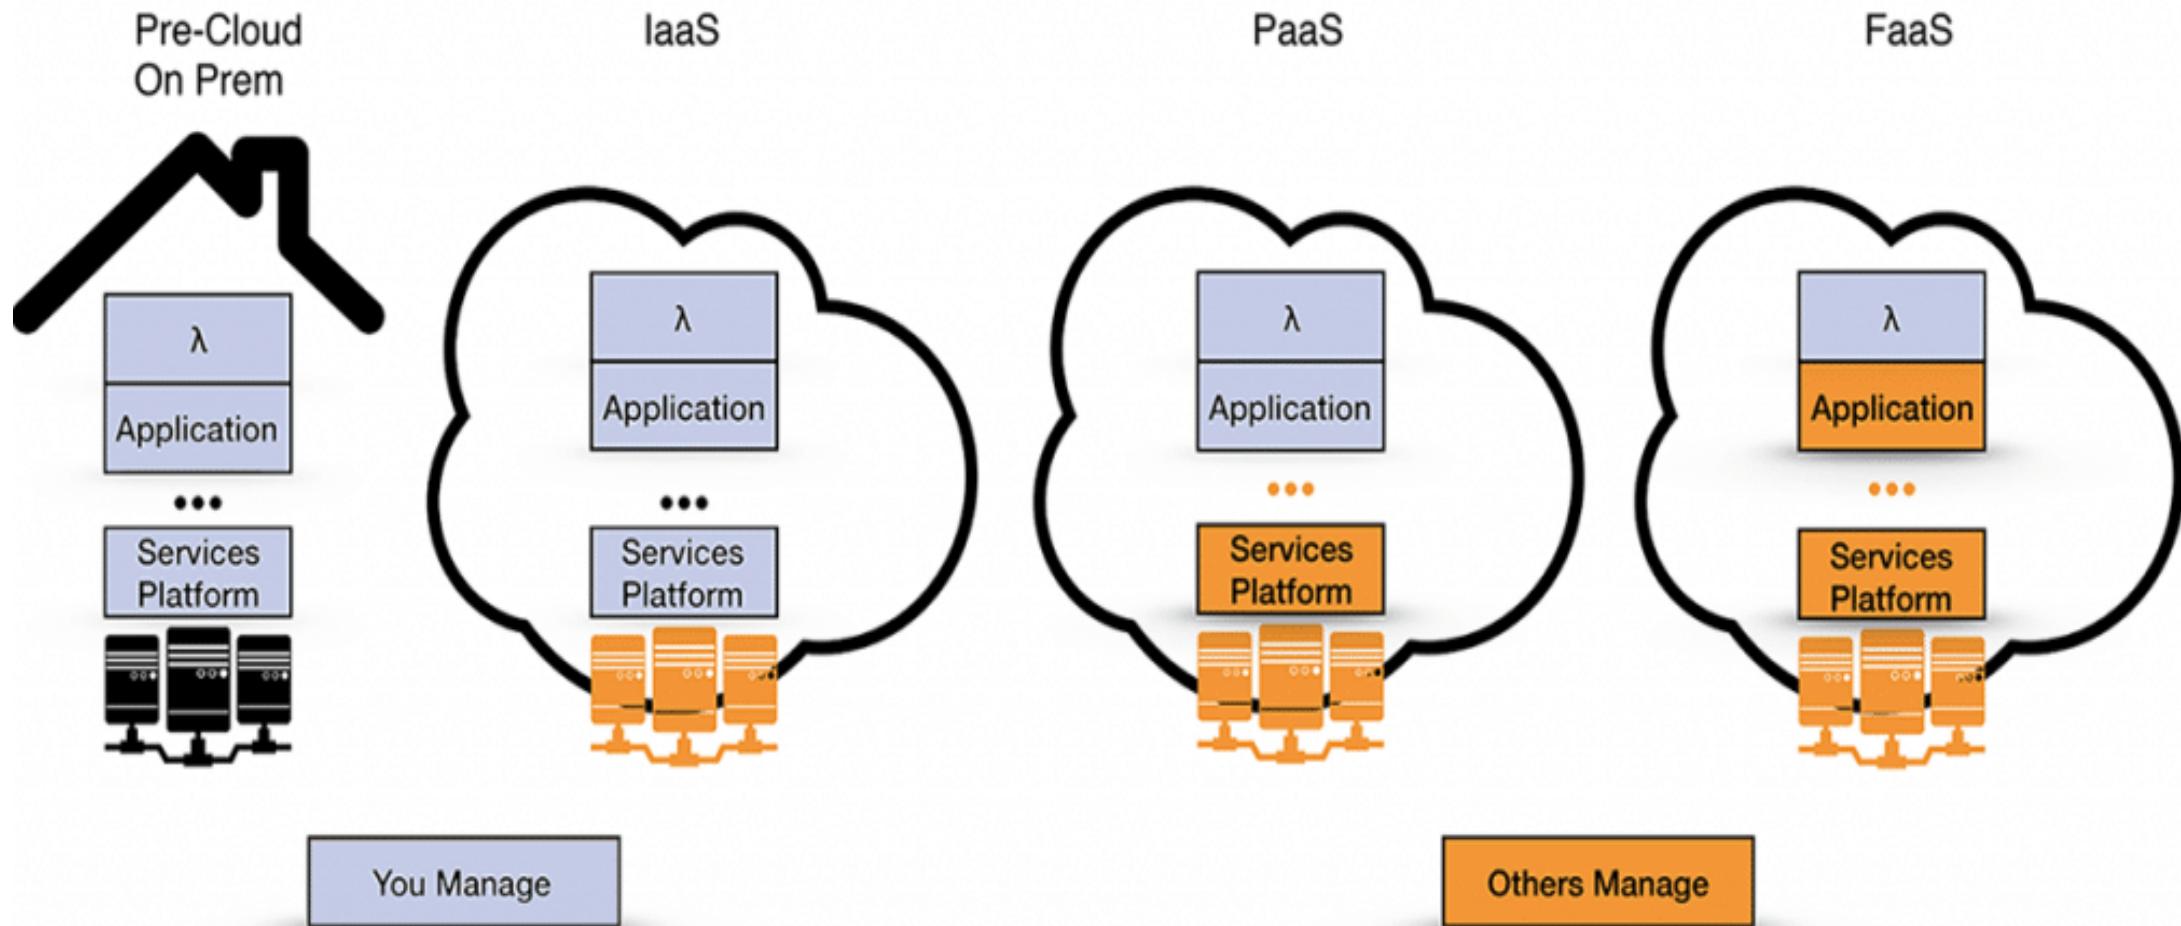
FaaS

- FaaS is a subset of serverless.
- FaaS, or Function-as-a-Service, is a cloud-computing service that allows customers to execute code in response to events, without managing the complex infrastructure typically associated with building and launching microservices applications.
- Hosting a software application on the internet typically requires provisioning and managing a virtual or physical server and managing an operating system and web server hosting processes.
- With FaaS, the physical hardware, virtual machine operating system, and web server software management are all handled automatically by the cloud service provider. This allows developers to focus solely on individual functions in their application code.

Advantages of FAAS

- Pay for execution
- Auto Scalable
- Faster Time to Market
- Polyglot Environment (Support for multiple programming languages)
- Highly Available

Evolution of Functions as a Service



FaaS vs. serverless

- Serverless and Functions-as-a-Service (FaaS) are often conflated with one another but the truth is that FaaS is actually a subset of serverless.
- Serverless is focused on any service category, be it compute, storage, database, messaging, api gateways, etc. where configuration, management, and billing of servers are invisible to the end user.
- FaaS, on the other hand, while perhaps the most central technology in serverless architectures, is focused on the event-driven computing paradigm wherein application code, or containers, only run in response to events or requests.

FaaS/Serverless computing providers

1. **Amazon Web Services (AWS) Lambda** is a popular serverless computing platform. It allows developers to write functions in various programming languages, including **Node.js, Python, Java, and C#**. These functions can be triggered by events such as HTTP requests, database updates, file uploads, or scheduled events.
2. **Azure Functions** is Microsoft's serverless computing offering. It supports languages like **C#, JavaScript, PowerShell, Python, and TypeScript**. Developers can write functions that respond to various triggers, including HTTP requests, timers, and messages from Azure services like Event Grid or Service Bus. Azure Functions automatically scales the resources to handle the workload, and users are billed based on the number of executions and the execution time.
3. **Google Cloud Functions** is a serverless computing platform provided by Google Cloud Platform (GCP). It supports languages such as **Node.js, Python, and Go**. Functions can be triggered by events from various sources like HTTP requests, Pub/Sub messages, or changes in Cloud Storage or Firestore. Google Cloud Functions scales automatically and charges users based on the number of function invocations, execution time, and memory consumption.
4. **IBM Cloud Functions** is the serverless computing offering from IBM Cloud. It supports multiple programming languages, including **Node.js, Swift, Python, and Java**. Users can create functions that respond to events from various sources like HTTP requests, database changes, or message queues. IBM Cloud Functions provides autoscaling capabilities and pricing is based on the number of function invocations and the execution time.

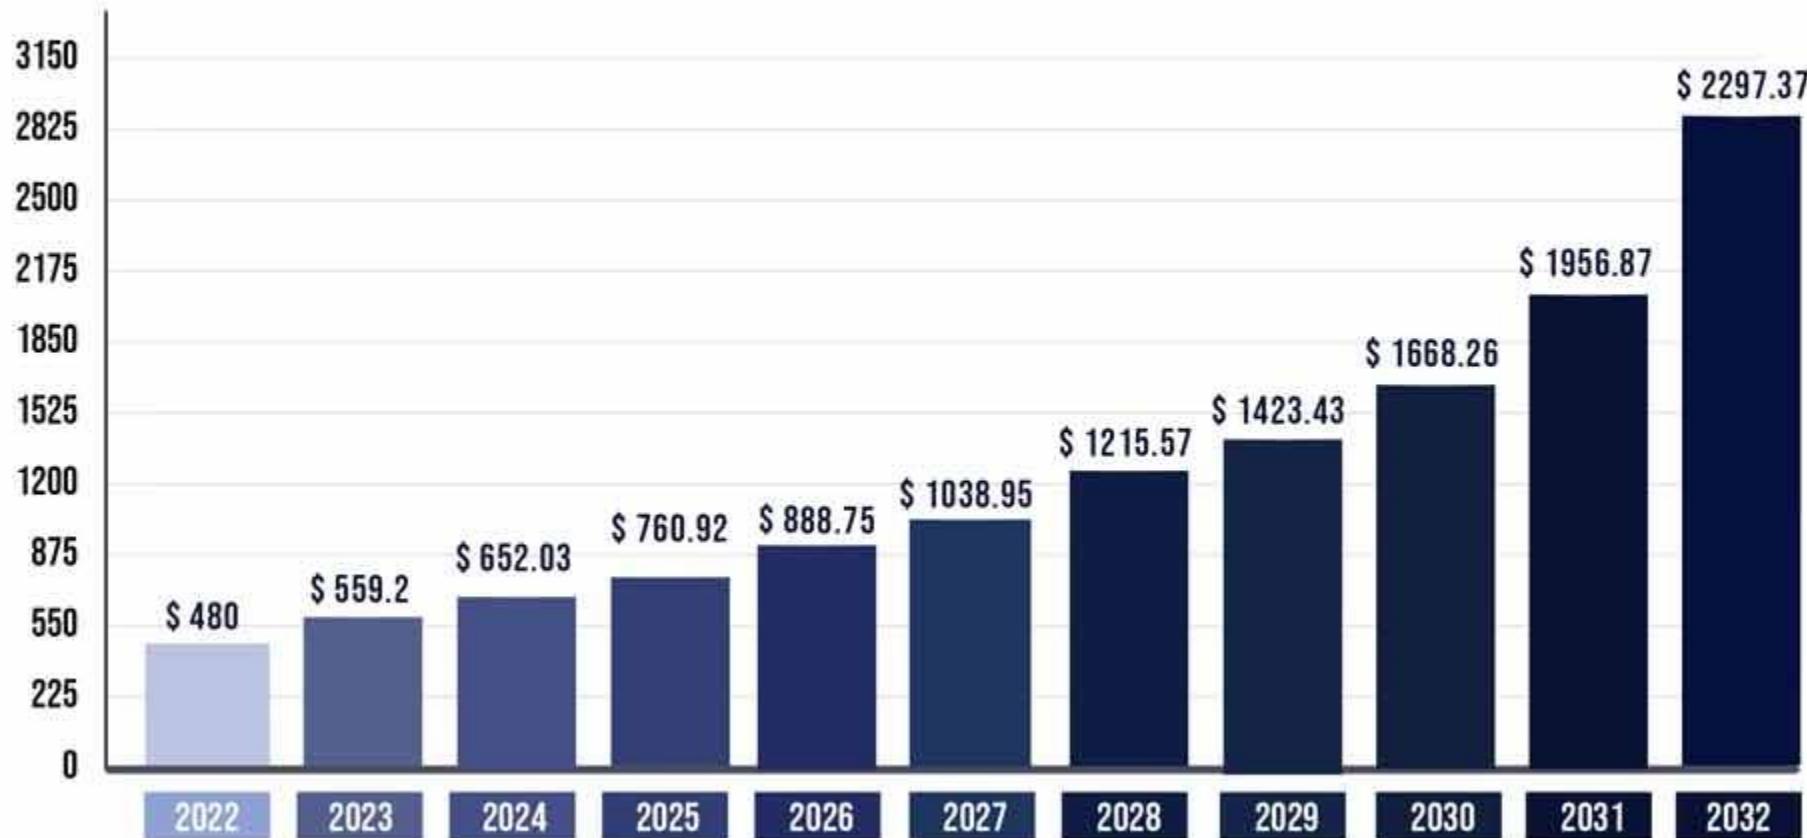
Cloud Computing Characteristics and Benefits

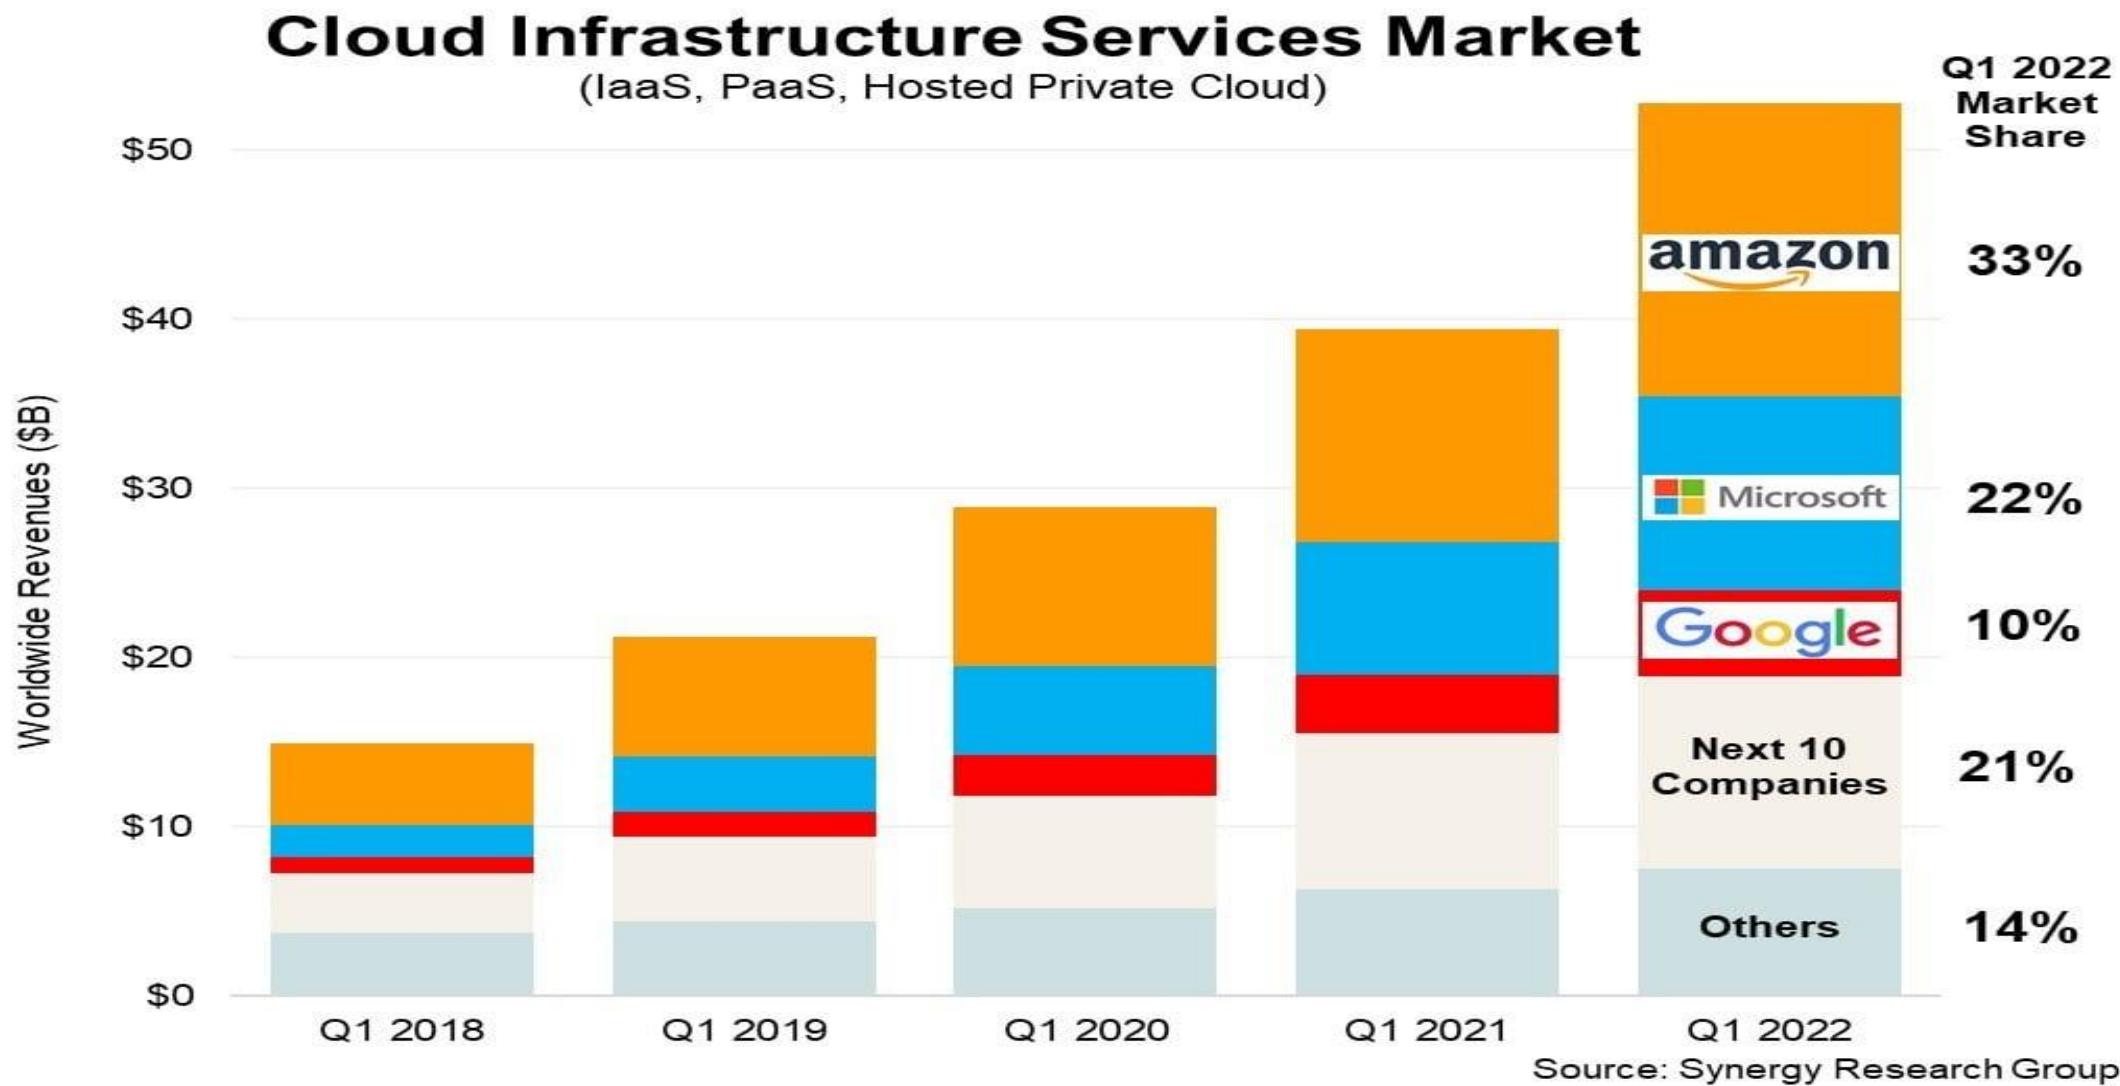
- No Upfront Commitments
- On demand access
- Nice pricing
- Simplified application acceleration and scalability
- Efficient resource allocation
- Energy efficiency and seamless creation and use third-party services.

Challenges Ahead

- Dynamic Provisioning of Cloud Computing Services
- Security and Privacy
- Legal issues
- Performance and Bandwidth Cost
- Reliability and Availability

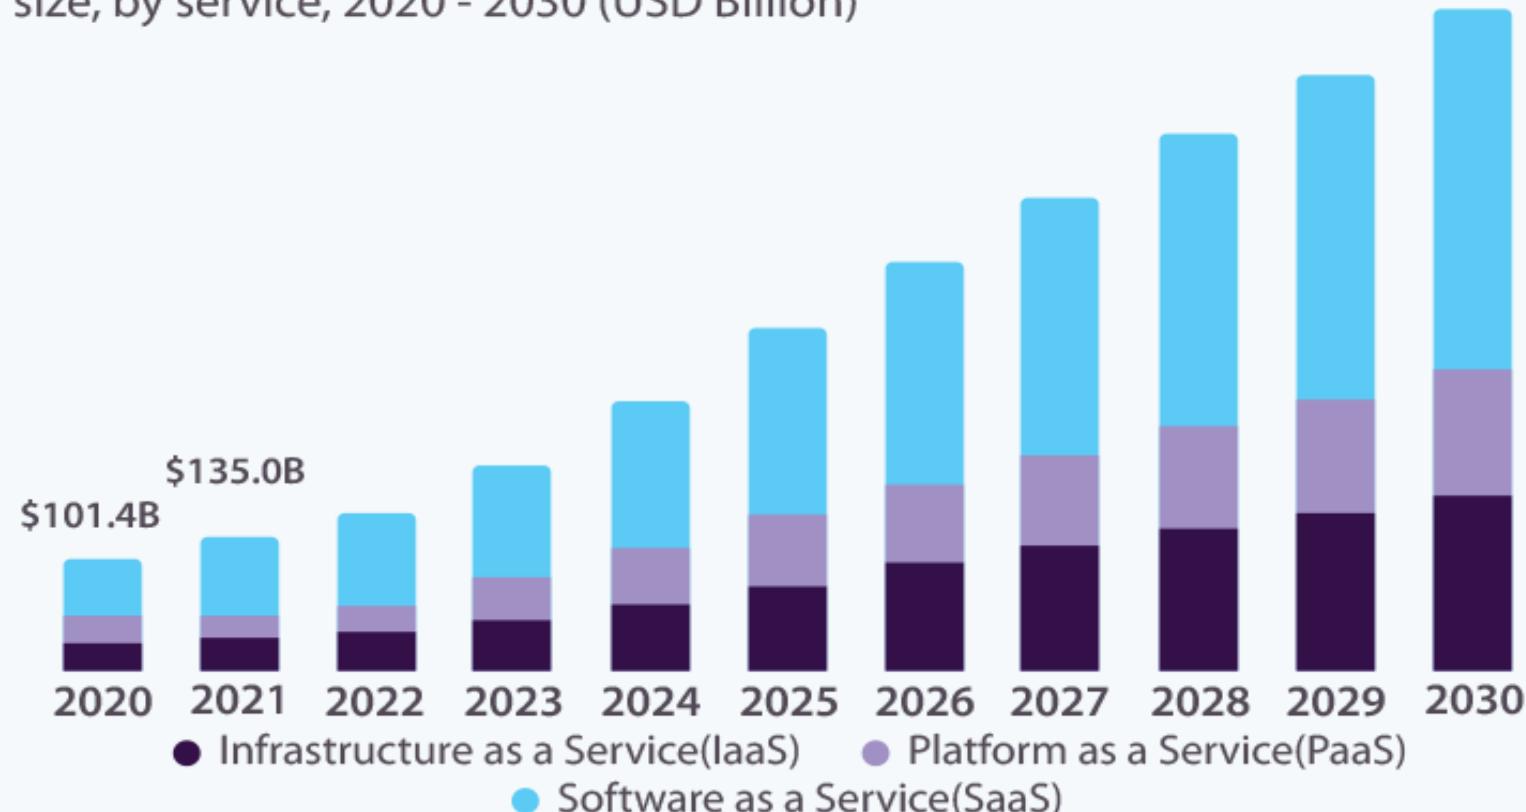
CLOUD COMPUTING MARKET 2022 TO 2032 [USD BILLION]

Source: www.precedenceresearch.com



U.S. Cloud Computing Market

size, by service, 2020 - 2030 (USD Billion)

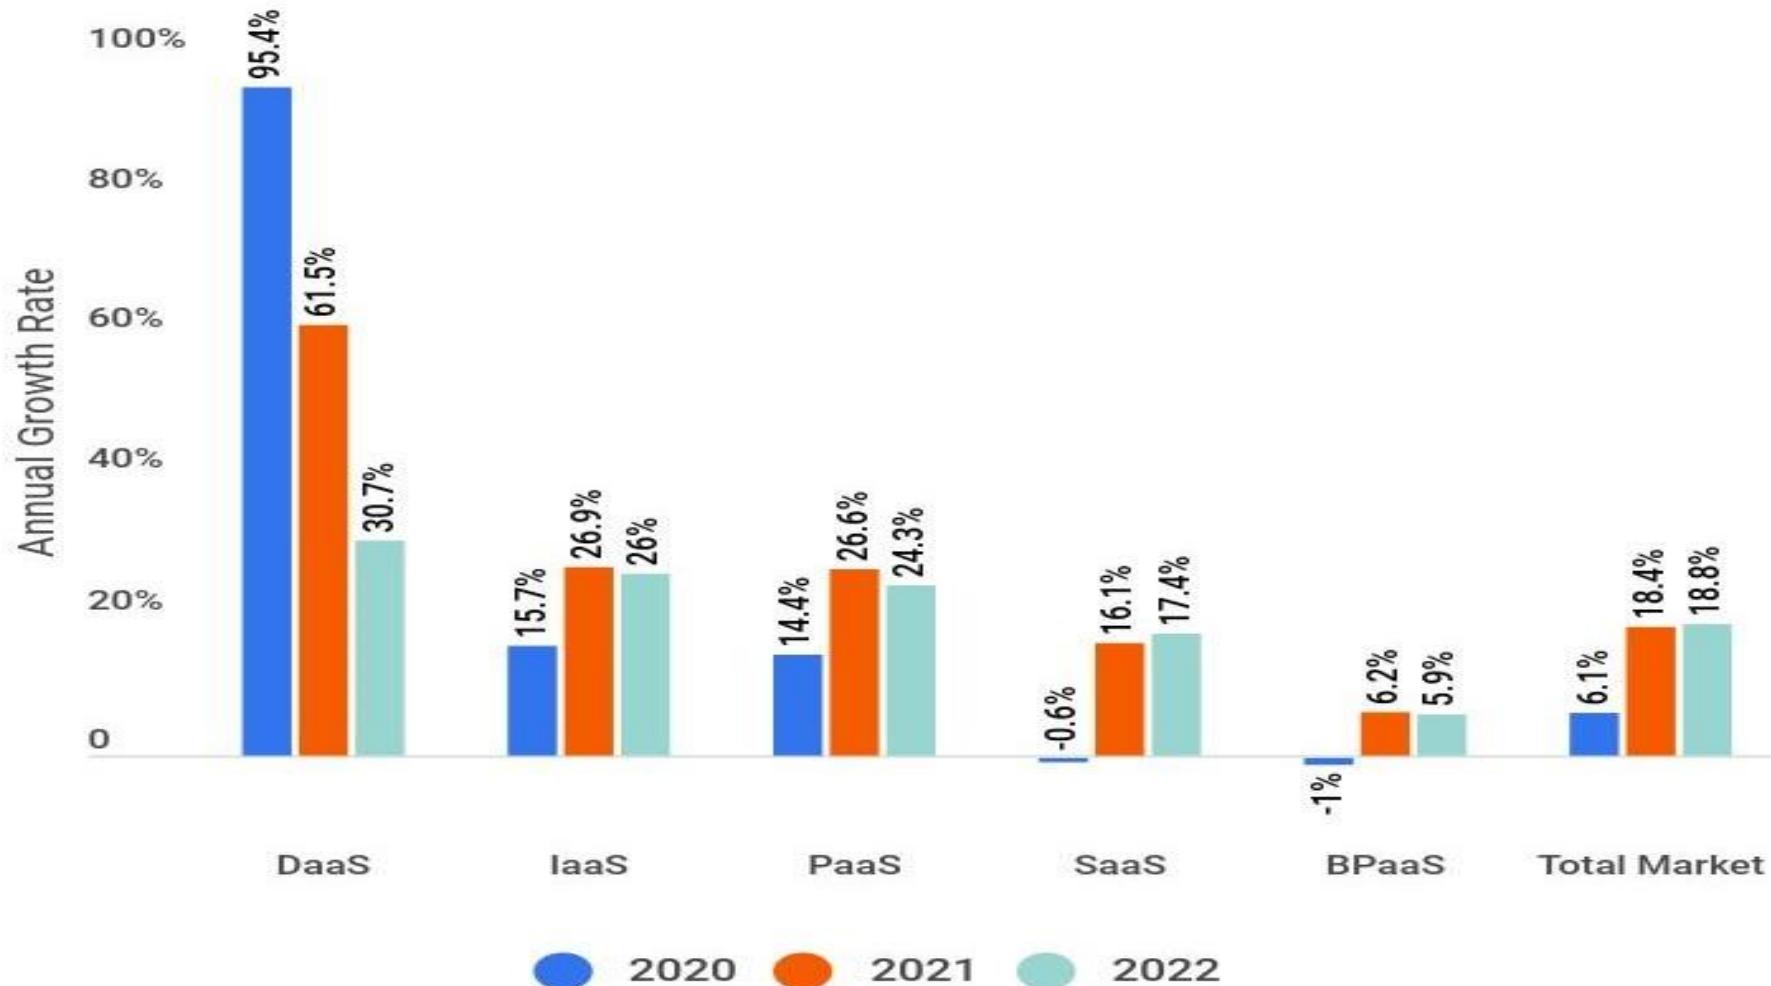


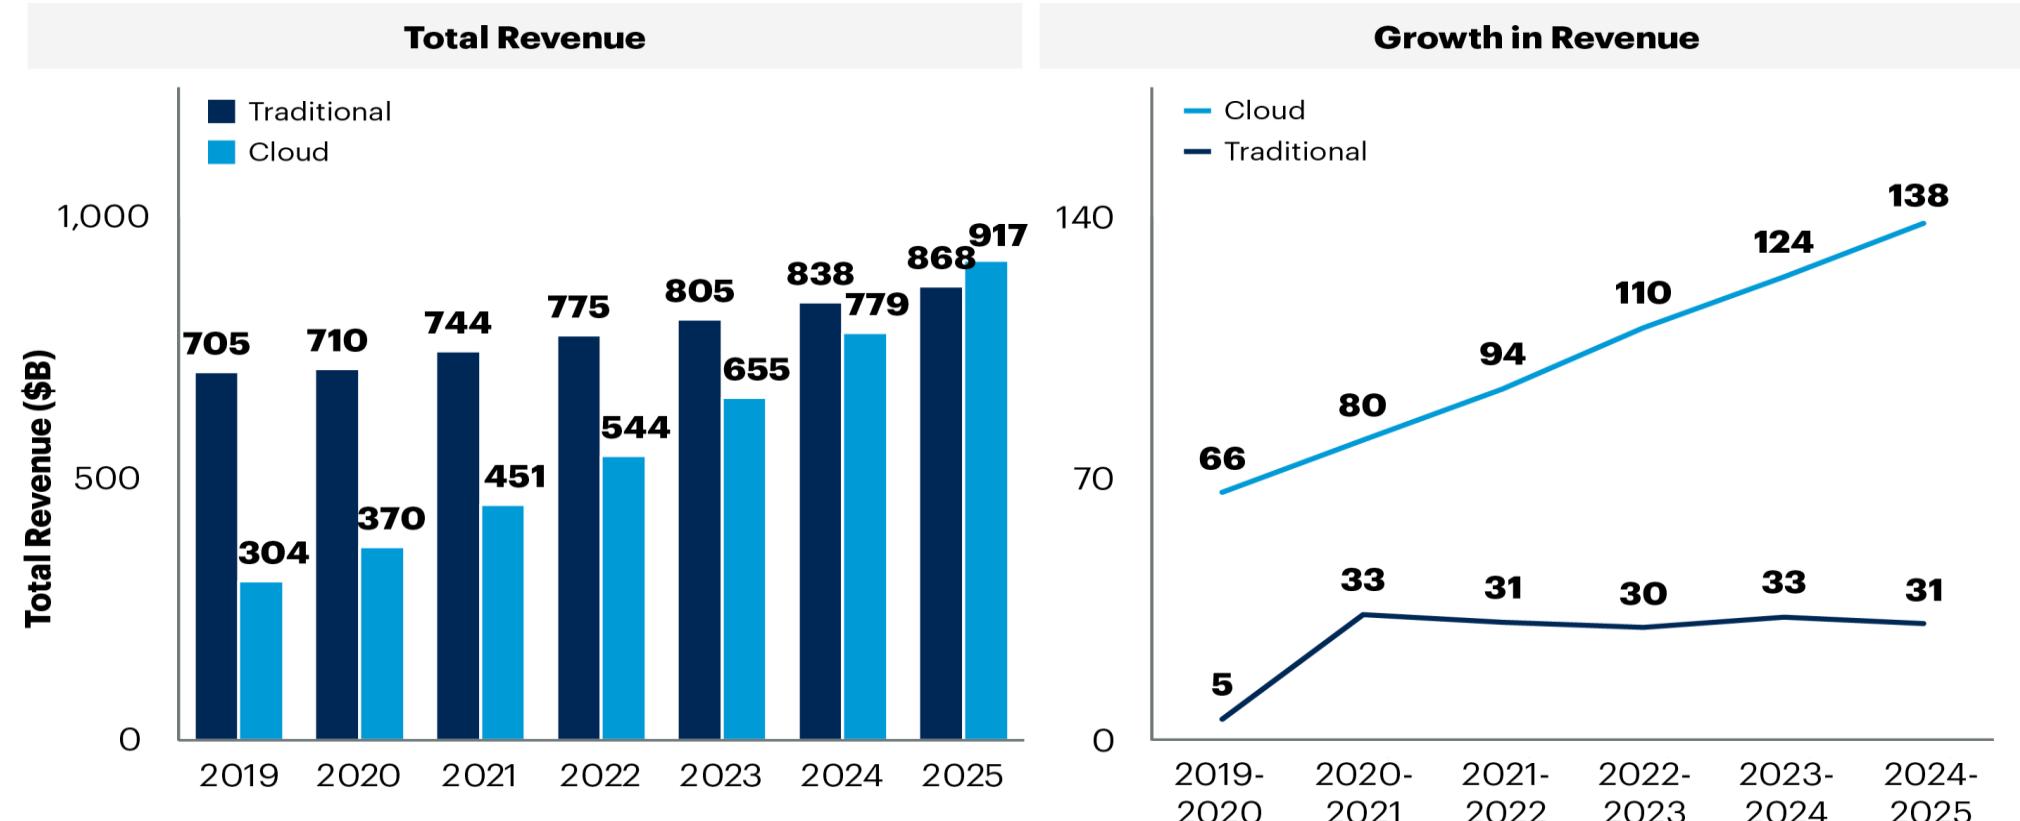
14.8%

U.S. Market CAGR,
2022 - 2030

Source:
www.grandviewresearch.com

PUBLIC CLOUD SERVICES ANNUAL GROWTH RATES

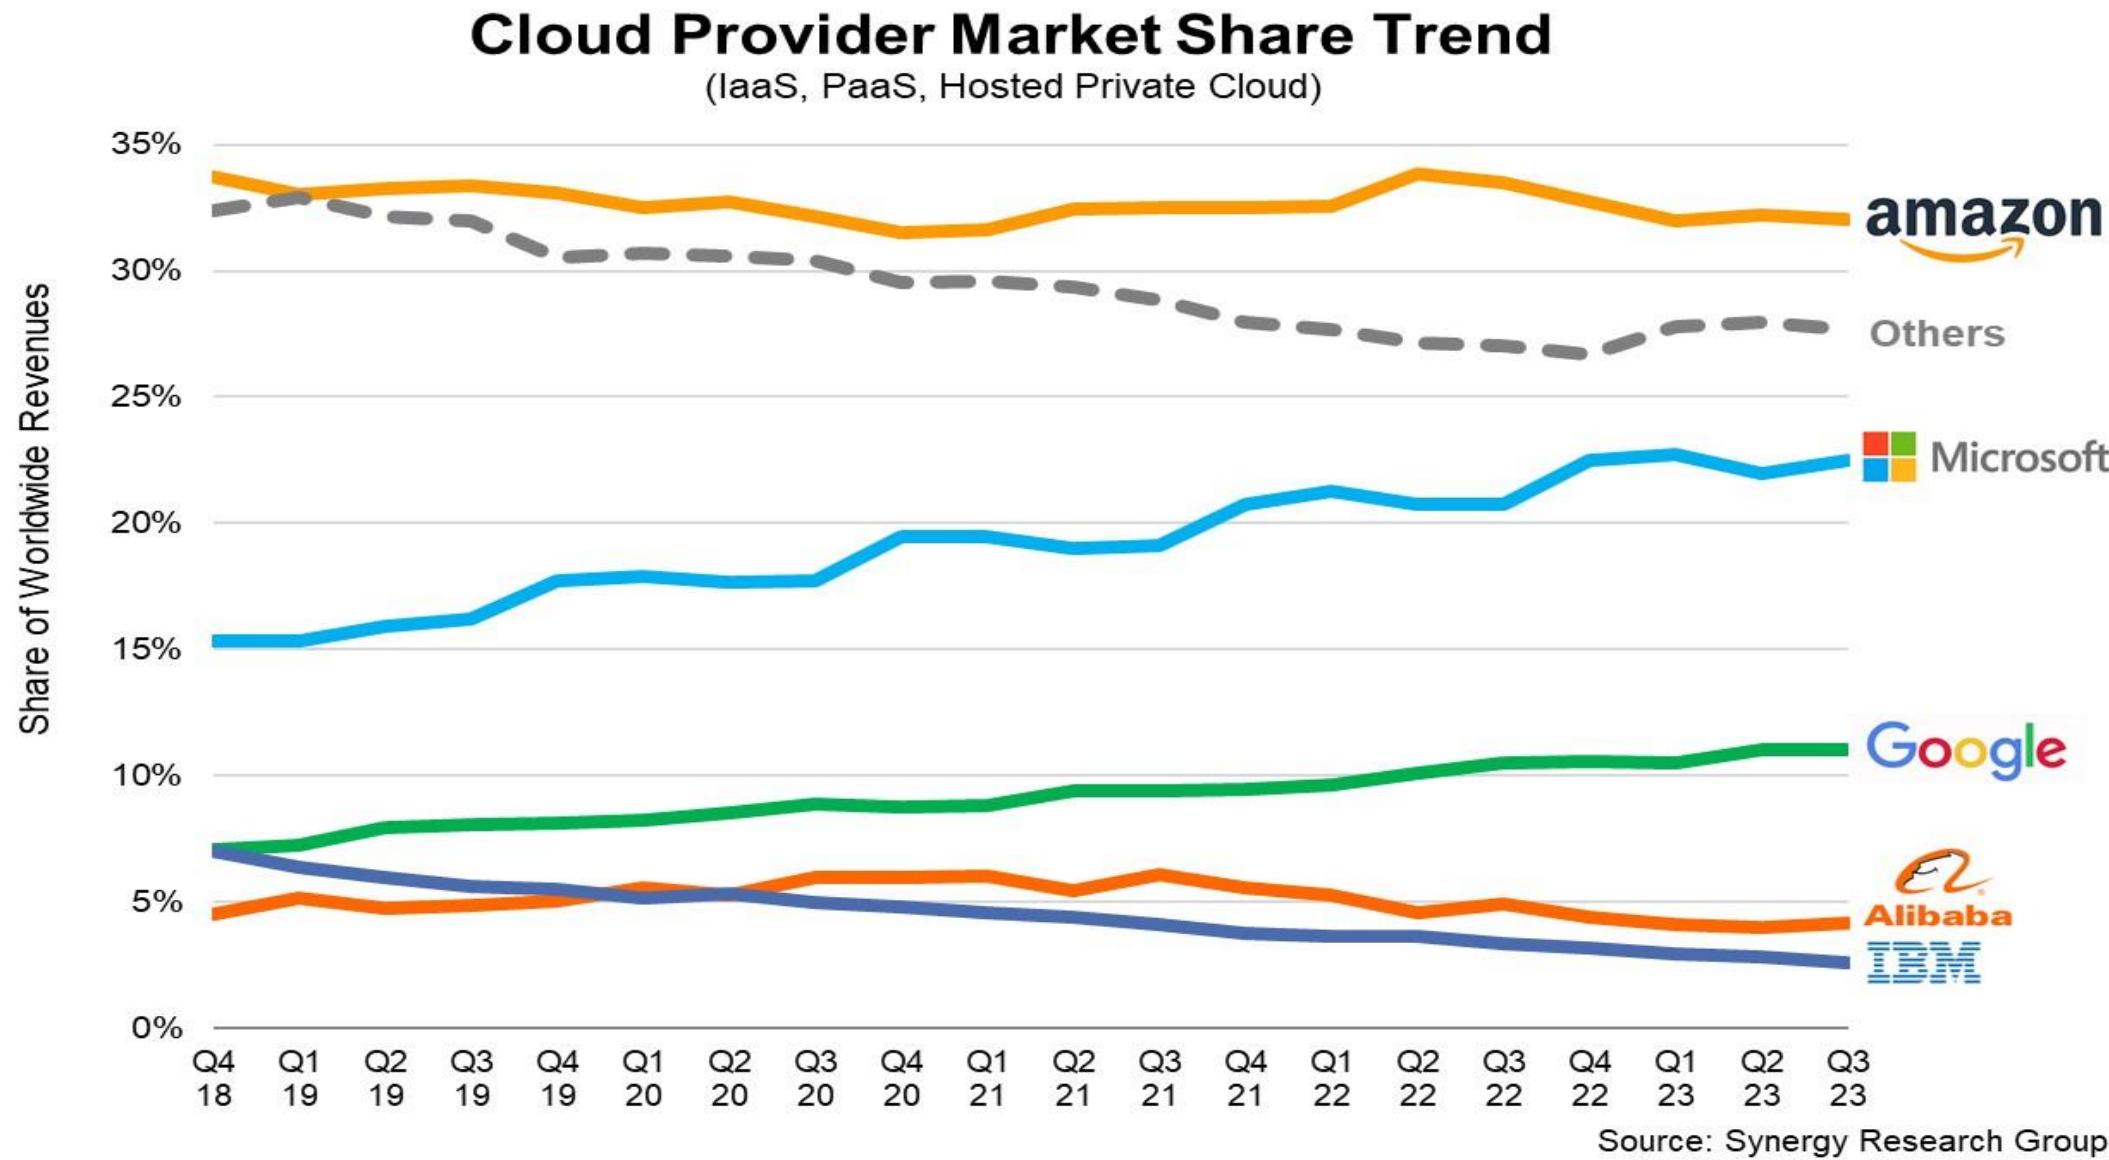




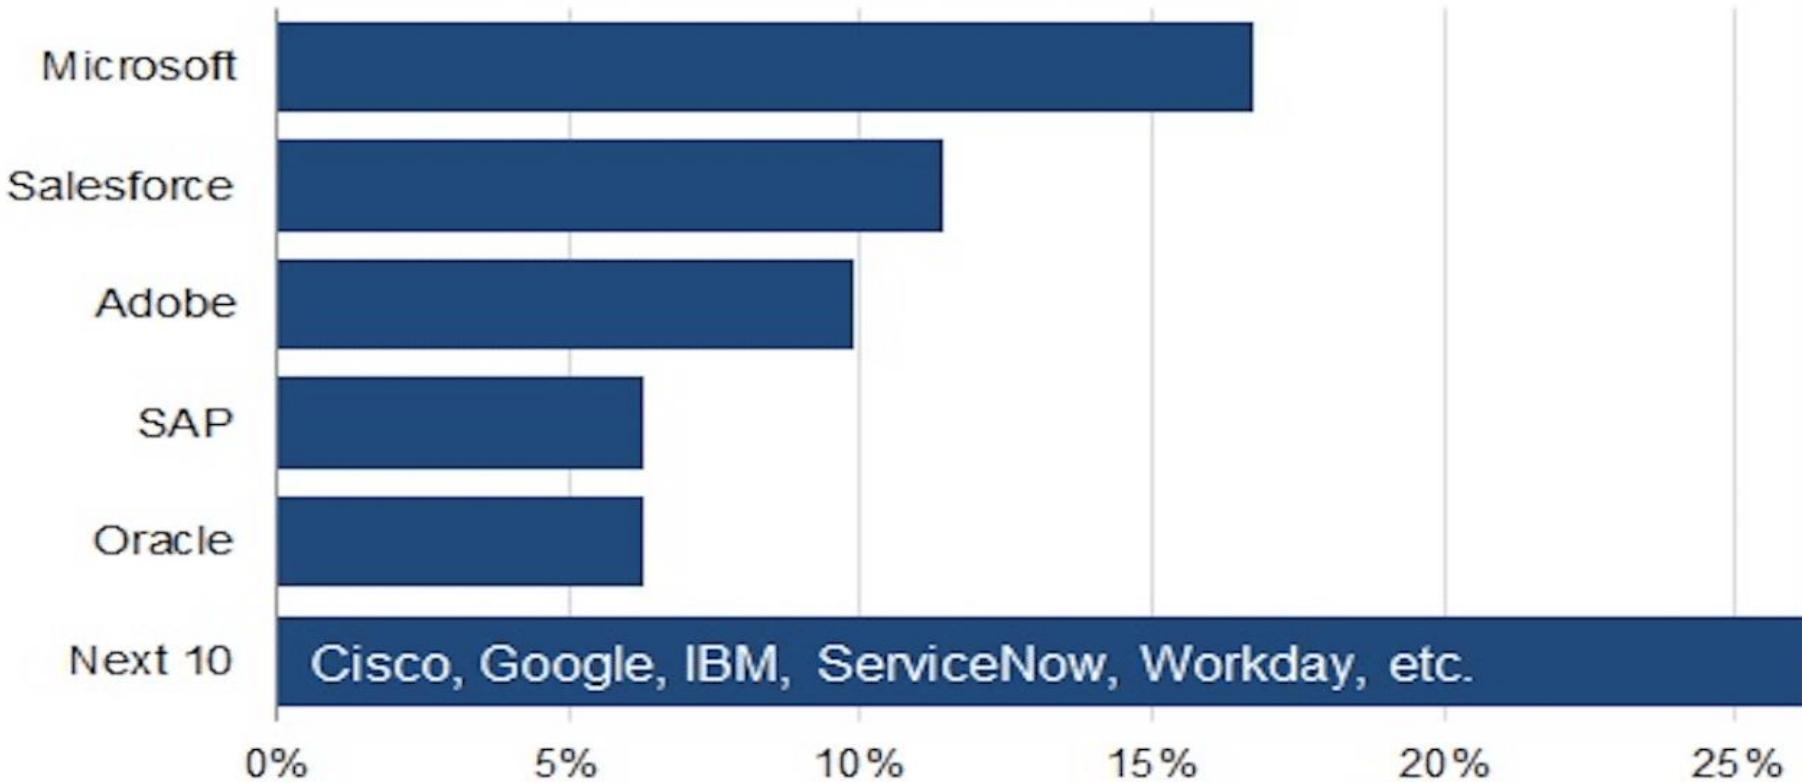
Source: Gartner

758067_C

Gartner®

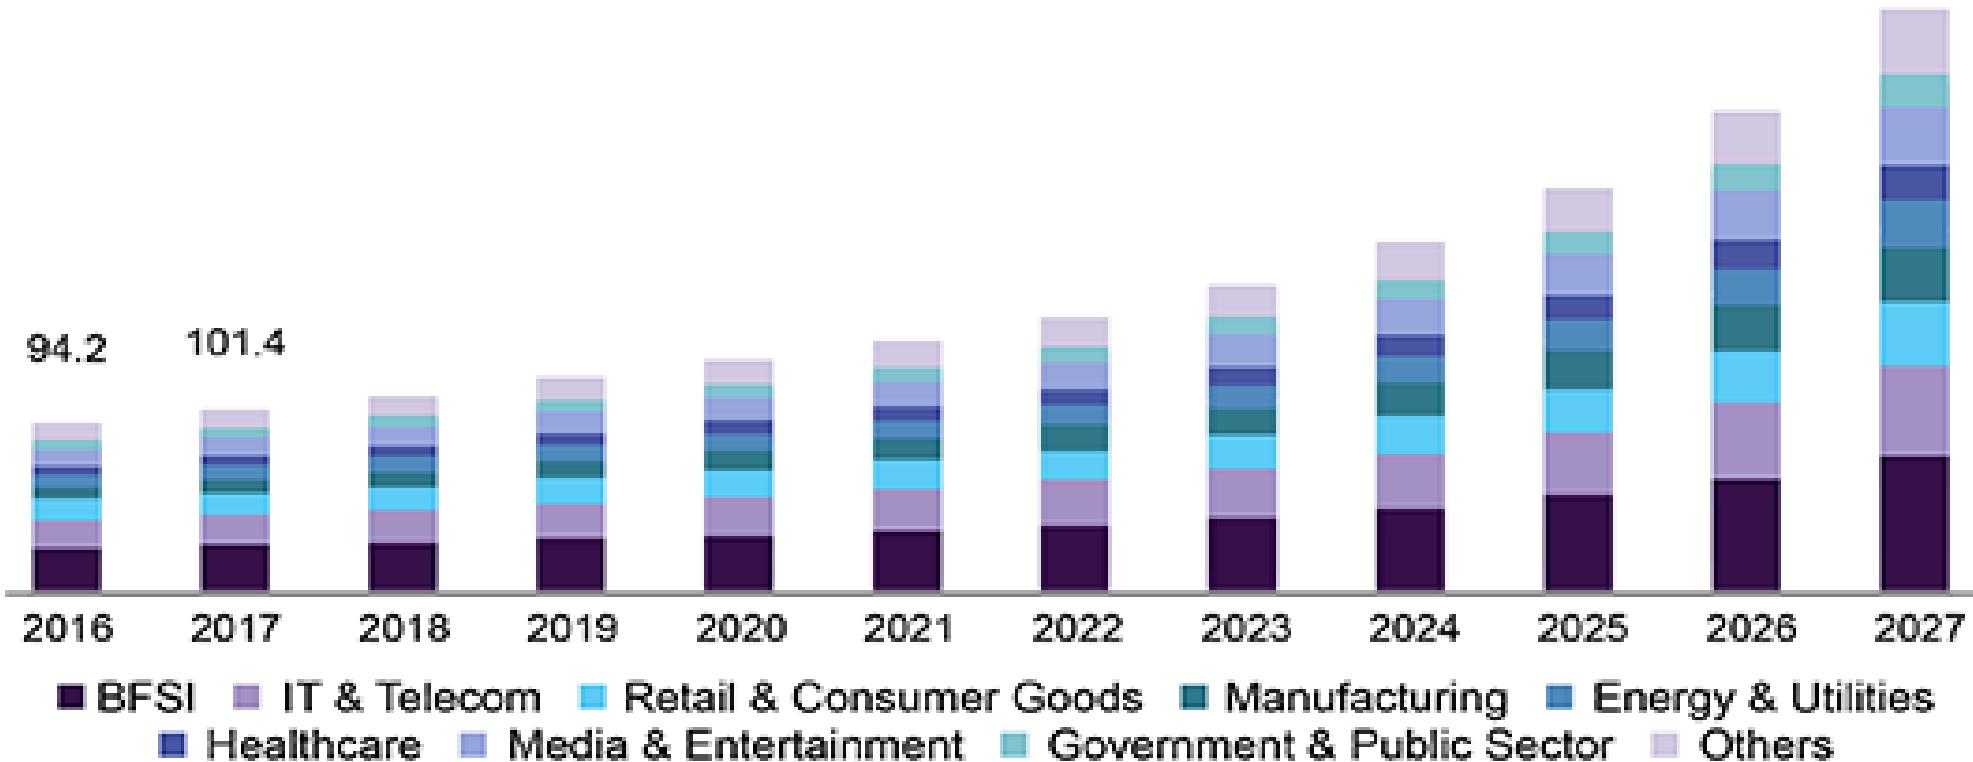


Enterprise SaaS Vendor Market Share & Revenue Growth





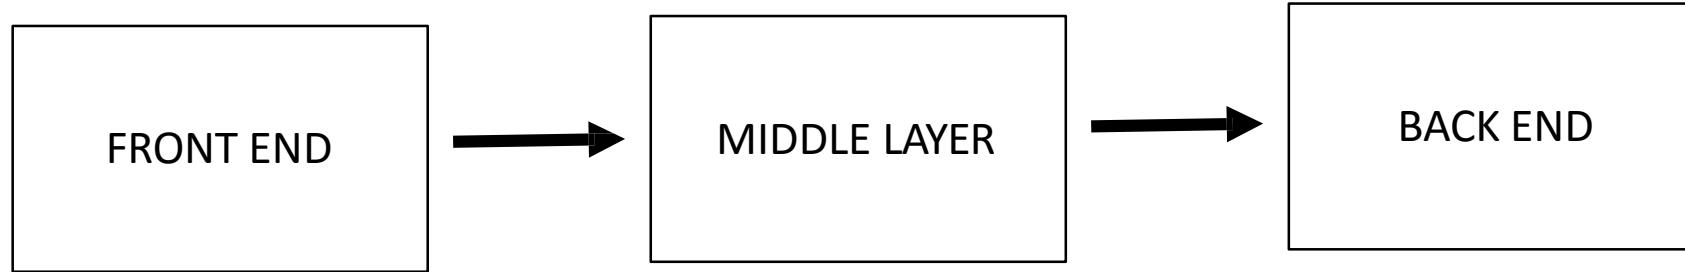
U.S. cloud computing market size, by end use, 2016 - 2027 (USD Billion)



Source: www.grandviewresearch.com

HOW CLOUD WORKS

- THREE TIER CLOUD SYSTEM



FRONT END:

- Is the part of the system that we see everyday.
- This interface could be in the form of applications installed on our devices or websites we retrieve through our web browsers.

HOW CLOUD WORKS

MIDDLE LAYER:

- Isn't necessarily physical. It is represented by software on both sides of the front end and the back end, and we call this middleware.
- Middleware is software that is used on both ends to ensure the front and the back can speak to each other.
- Administered by a central server and follows a specific sets of rules, known as protocols.

BACK END:

- It's where the actual cloud part of the system is housed.
- This is made of computers, servers, and databases and storage systems.

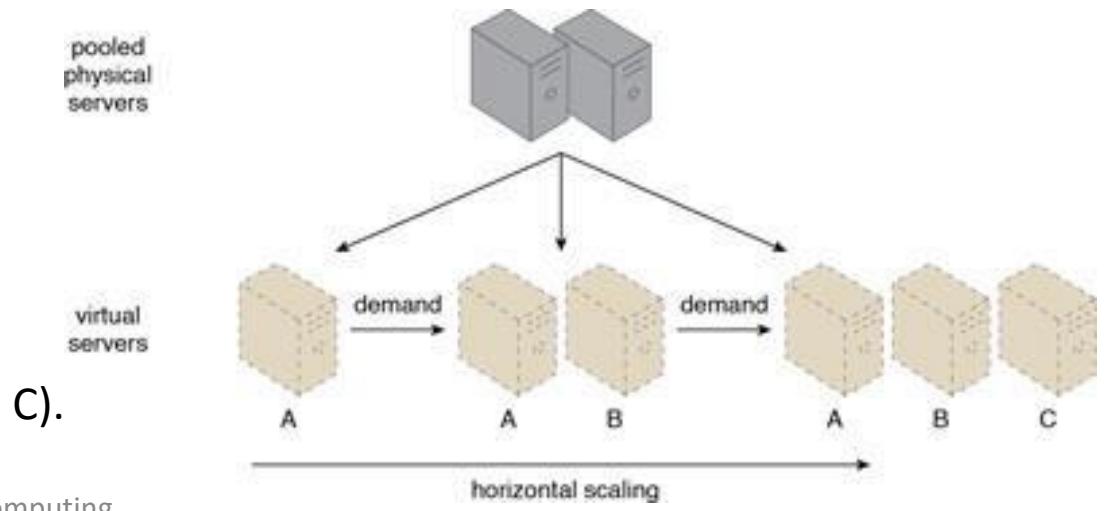
Types of scaling

The following are the types of scaling:

1. Horizontal Scaling-scaling out and scaling in.

The allocating or releasing of IT resources that are of the same type is referred to as horizontal scaling. The horizontal allocation of resources is referred to as scaling out and the horizontal releasing of resources is referred to as scaling in.

Figure 1 - An IT resource (Virtual Server A) is scaled out by adding more of the same IT resources (Virtual Servers B and C).



Types of scaling

2. Vertical Scaling- scaling up and scaling down

When an existing IT resource is replaced by another with higher or lower capacity, vertical scaling is considered to have occurred.

To replace an IT resource with another that has a higher is referred to as scaling up and the replacing an IT resource with another that has a lower capacity is considered scaling down.

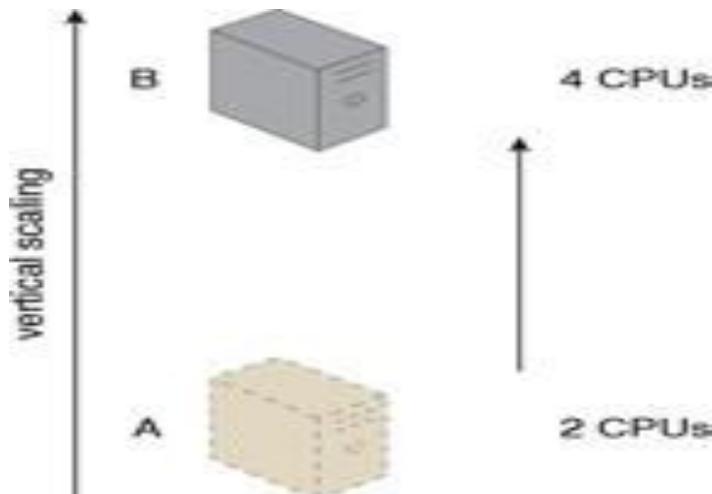


Figure 2 - An IT resource (a virtual server with two CPUs) is scaled up by replacing it with a more powerful IT resource with increased capacity for data storage (a physical server with four CPUs).

Horizontal vs Vertical Scaling

| Horizontal Scaling | Vertical Scaling |
|--|---|
| less expensive (through commodity hardware components) | more expensive (specialized servers) |
| IT resources instantly available | IT resources normally instantly available |
| resource replication and automated scaling | additional setup is normally needed |
| additional IT resources needed | no additional IT resources needed |
| not limited by hardware capacity | limited by maximum hardware capacity |

Table 1 - A comparison of horizontal and vertical scaling.

Evolution of cloud technologies

- *Distributed Systems*

- A distributed system is a collection of independent computers that appears to its users as a single system and also it acts as a single computer.
- The main and primary motive of distributed systems is to share resources and to utilize them better.

- This is absolutely true in case of cloud computing because in cloud computing we are sharing the single resource by paying rent.
- The resource is single because the definition of cloud computing clearly states that in cloud computing the single central copy of a particular software is stored in a sever (which is located on a anonymous location) and users are accessing that on PAY PER USE BASIS.

Characteristics of Distributed Applications

- Modern applications consume data from distributed data sources to fulfill users expectations
- Distributed applications are designed for:
 - Scalability
 - Low latency
 - Availability
 - Reliability
 - Security and privacy

Logical Layers of Distributed Applications

- Separation of concerns (Separating the responsibilities) between different components helps to achieve better maintainability, testability, and agility.
- It is easier to test each layer separately rather than testing the whole system together.
- Applications are designed in layers:
 1. Data Layer
 2. Business Layer
 3. User Interface Layer
 4. Service Layer

Data Layer

- The data layer is responsible for storing, querying, updating, or deleting the data as required while maintaining a reasonable performance.
- This can be a complicated task when you are dealing with a large set of data, distributed across several data sources.
- Data can be replicated, distributed, and handled according to its characteristics. For example, client contacts can be replicated across the data center because they change slowly. However, information about stocks must be always accurate and therefore must be read from a single source.

Business Layer

- The business layer(execution layer) contains the business logic and is responsible for carrying out the use-case scenarios of the application.
- This layer implements the logic of the application.
- The business logic uses the data layer to read and store data, and the UI layer to interact with the client.
- The execution layer contains all the algorithms and logic of the application and is considered the brain of the application.

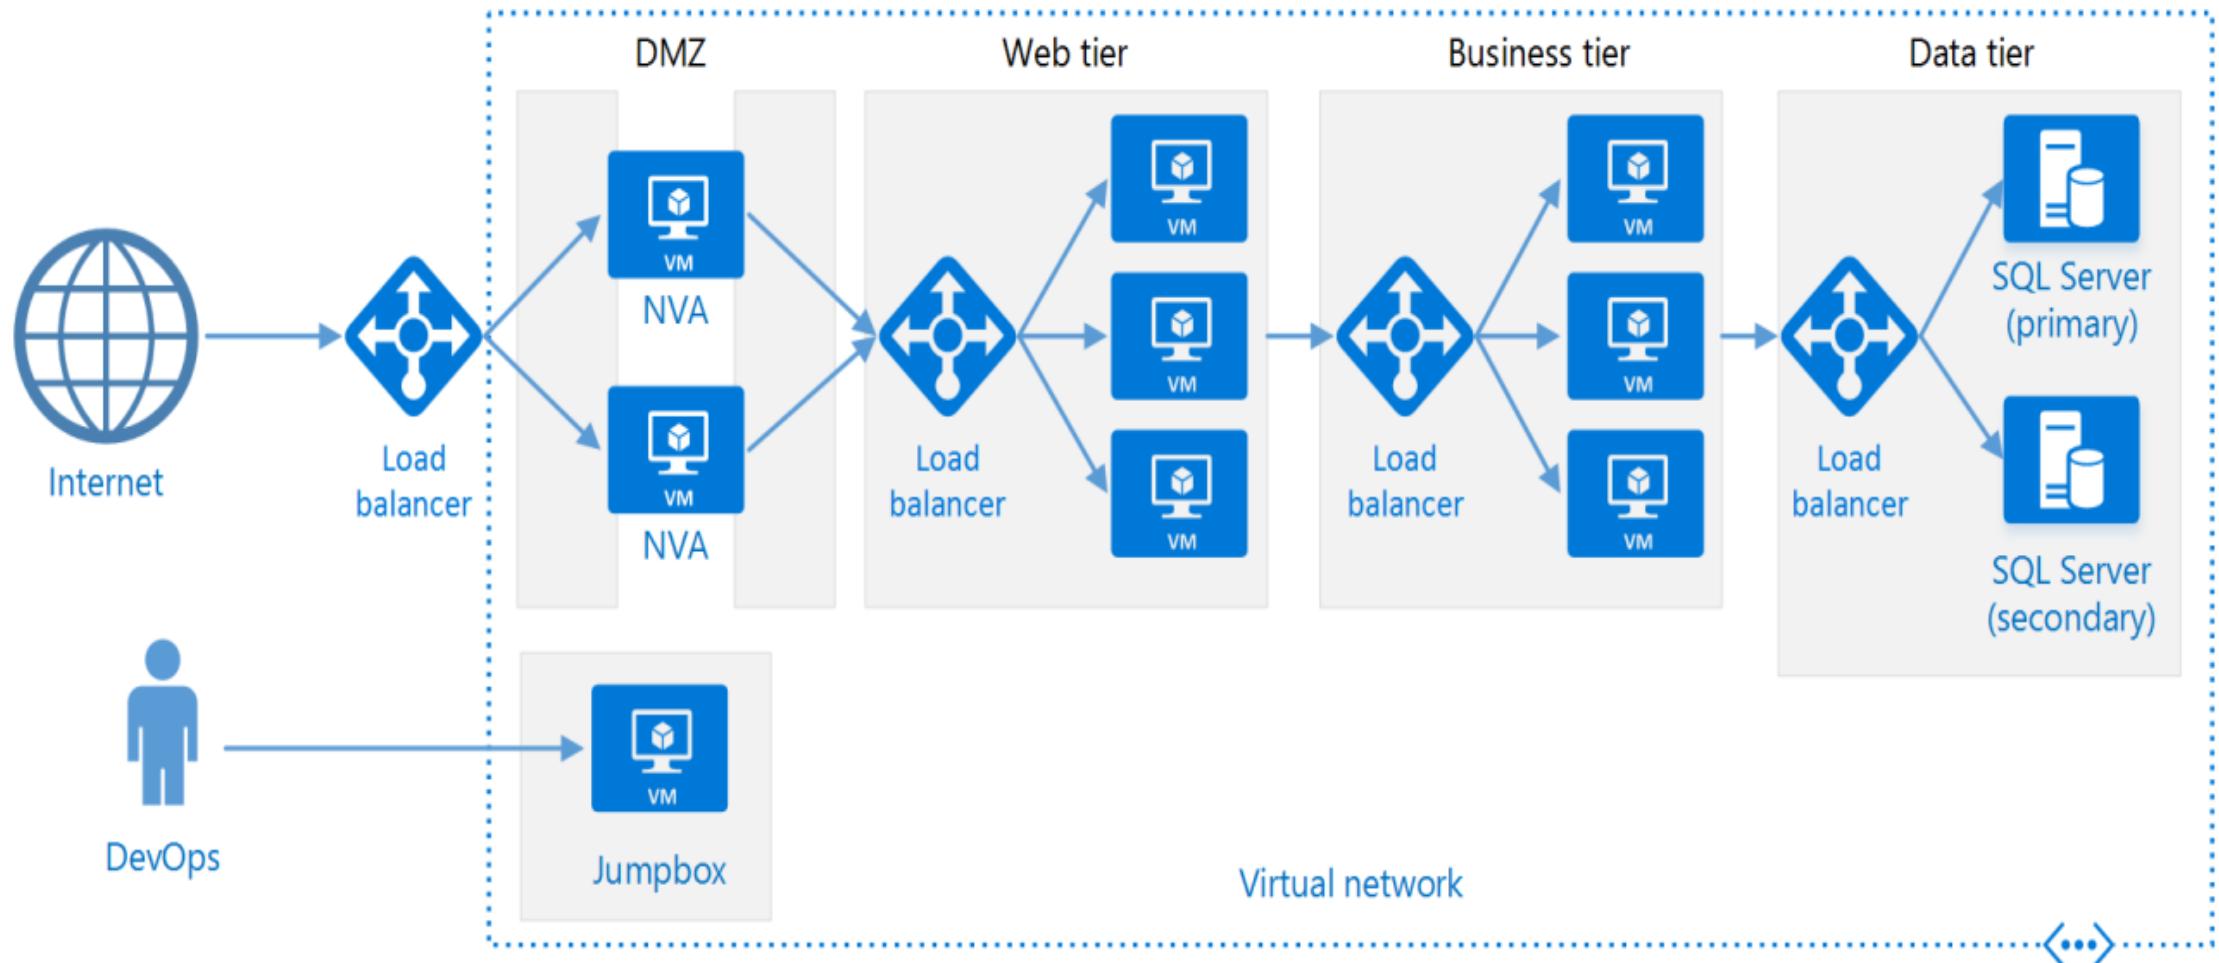
Service Layer

- The service layer exposes some of the capabilities of the application to the world as services.
- Other applications might consume these services and use them as a data source or as a remote execution engine.
- The service layer acts as the interface for other applications, in contrast to the user interface layer, which targets humans.
- The service layer drives collaboration of applications and enables distribution of computing load and data.
- It is responsible for defining a contract that consumers must maintain to use the service.
- It enforces security policies, validates incoming requests, and maintains the application resources.

User Interface Layer

- The user interface layer is the layer through which users interact with the application.
- It visually depicts the data and operations of the application effectively and provides the users a simplified medium for consuming the application data.
- While designing the UI, developers must consider the varying expectations of different people and cultures.
- The UI should always be responsive, yet its ability to respond quickly might be CPU-intensive especially when using modern interfaces such as touchscreens.
- The UI can be displayed on a variety of different devices, some of which might have extreme limitations such as screen size and resolution.
- Nevertheless, the UI must be effective and present a useful visualization.
- The UI must provide simple, yet effective methods for the user to enter data and activate the business and data layers to store and process it. Proper UI design is crucial because if the UI is not user friendly, the application will not be used.

Three tiers Applications in Cloud



Mainframes

- A large high-speed computer, especially one supporting numerous workstations or peripherals.
- The central processing unit and primary memory of a computer.



Clusters

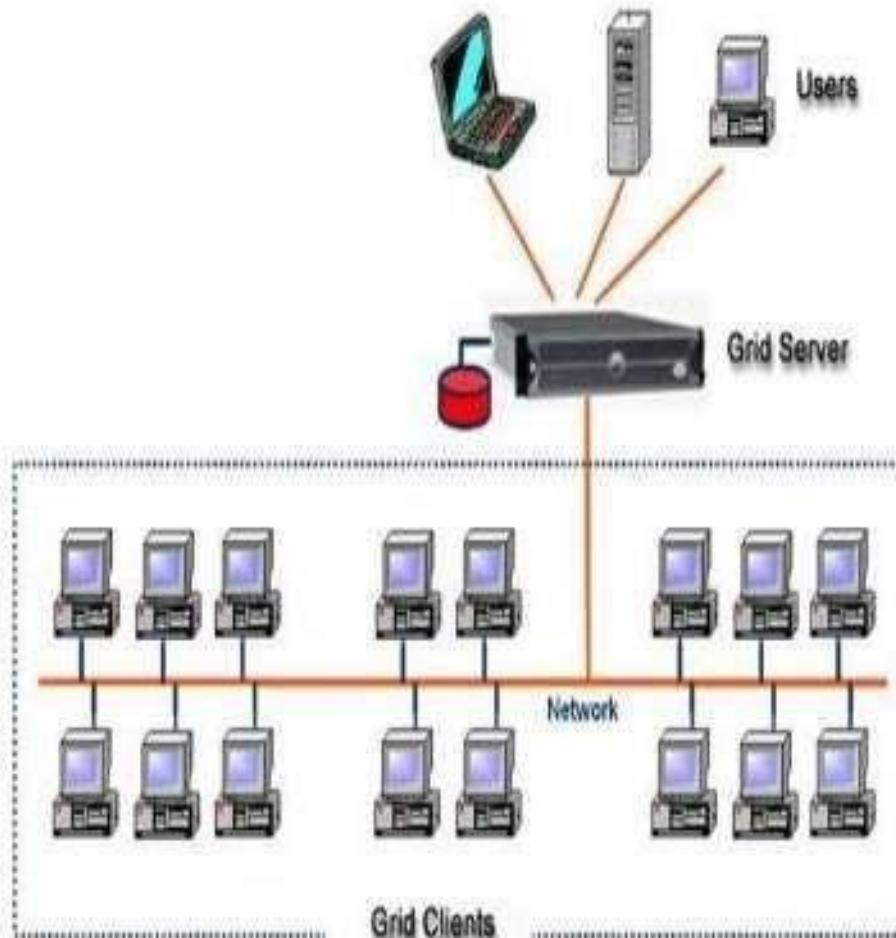
- A **computer cluster** consists of a set of loosely or tightly connected computers that work together so that, in many respects, they can be viewed as a single system.
- computer clusters have each node set to perform the same task, controlled and scheduled by software.



Grids

- **Grid computing** is the collection of **computer** resources from multiple locations to reach a common goal. The **grid** can be thought of as a **distributed** system with non-interactive workloads that involve a large number of files.
- The components of a cluster are usually connected to each other through fast local area networks, with each *node* (computer used as a server) running its own instance of an operating system. In most circumstances,

How Grid computing works ?

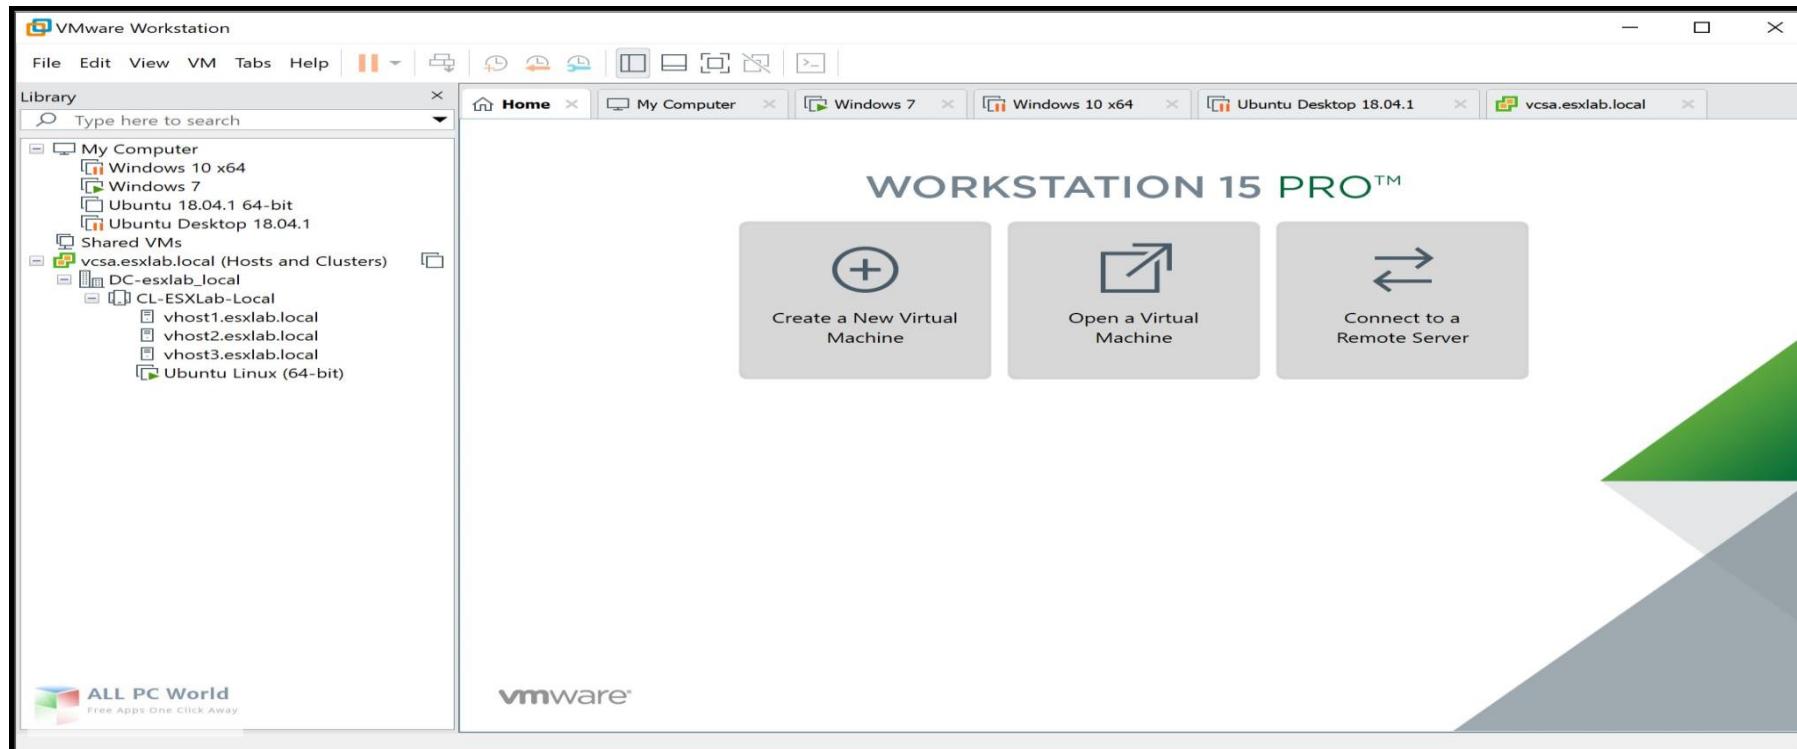


In general, a grid computing system requires:

- **At least one computer, usually a server, which handles all the administrative duties for the System**
- **A network of computers running special grid computing network software.**
- **A collection of computer software called middleware**

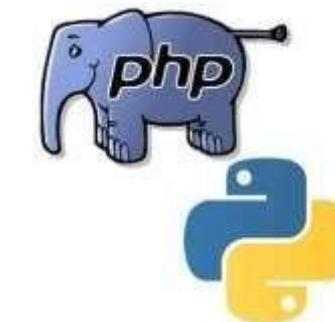
Virtualization

- In computing, **virtualization** refers to the act of creating a virtual (rather than actual) version of something, including virtual computer hardware platforms, storage devices, and computer network resources.
- Another Cloud Technology of Cloud Computing
- Hardware virtualization
- Storage or Network Virtualization



Web 2.0

- The second stage of development of the Internet, characterized especially by the change from static web pages to dynamic or user-generated content and the growth of social media.

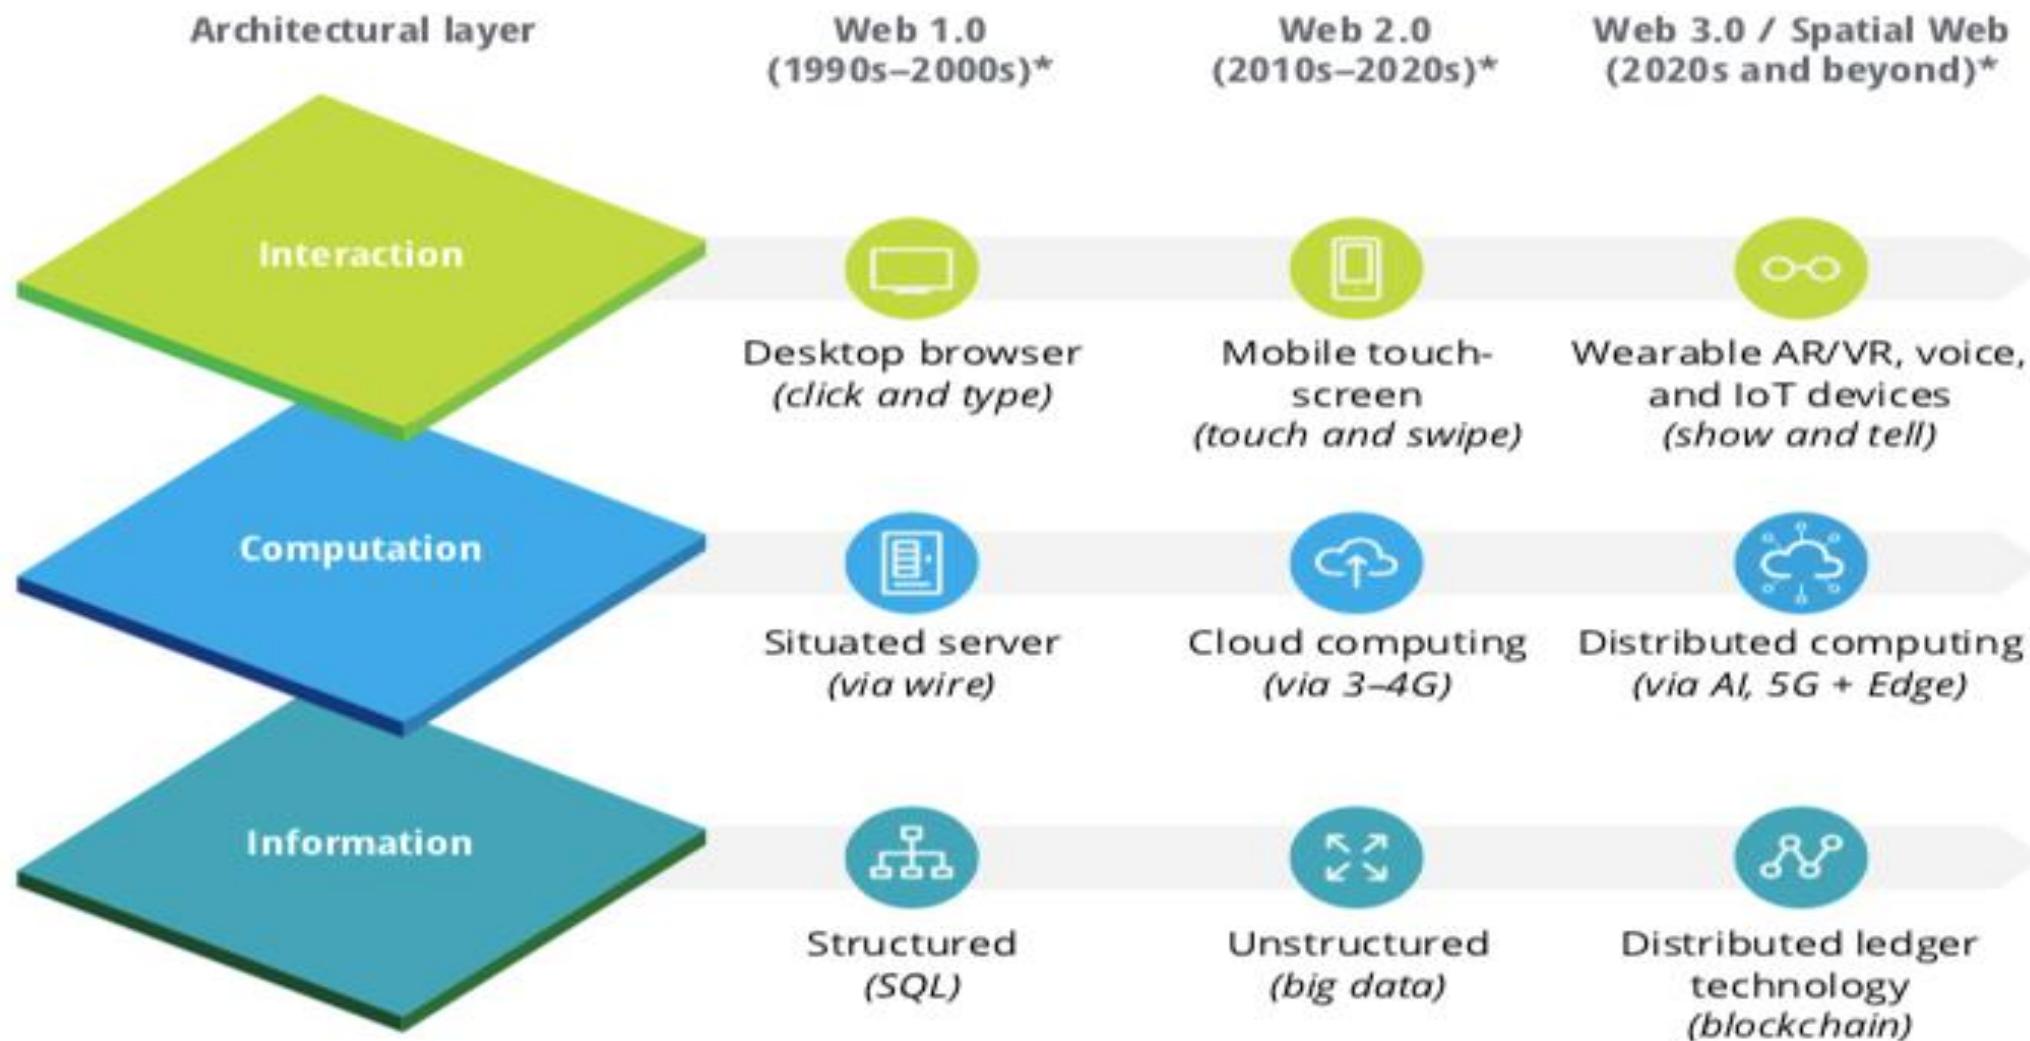




Cloud Computing
Spring 2024

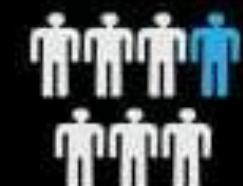
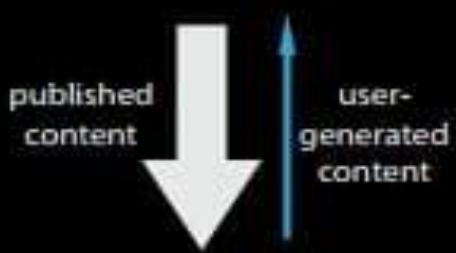
Web 3.0

- Web 3.0 also known as Web 3, is a concept for the next generation of the World Wide Web.
- The vision for Web 3.0 is that it will be a version of the internet that is both open and decentralized.
- Decentralization means that the internet isn't controlled by a single body like the government but by a community of members. It is millions of devices linked together in an open network.



Web 1.0

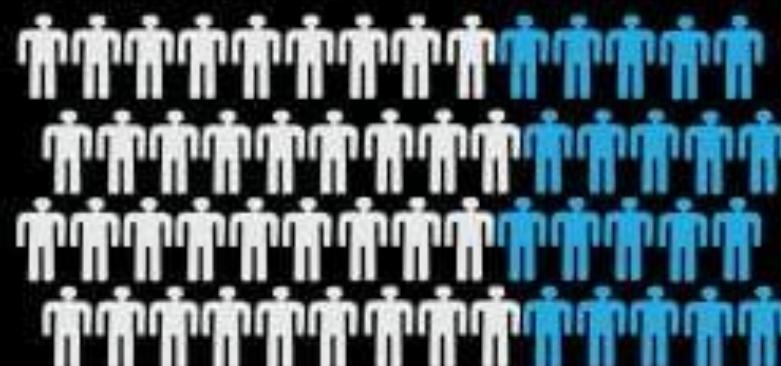
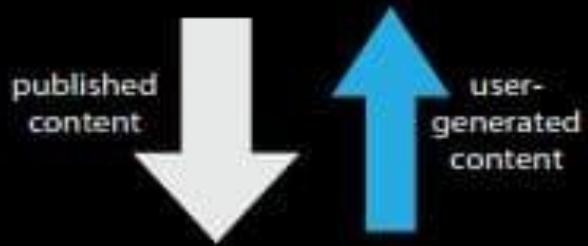
100,000 websites
(read-only Web)



50,000,000 users

Web 2.0

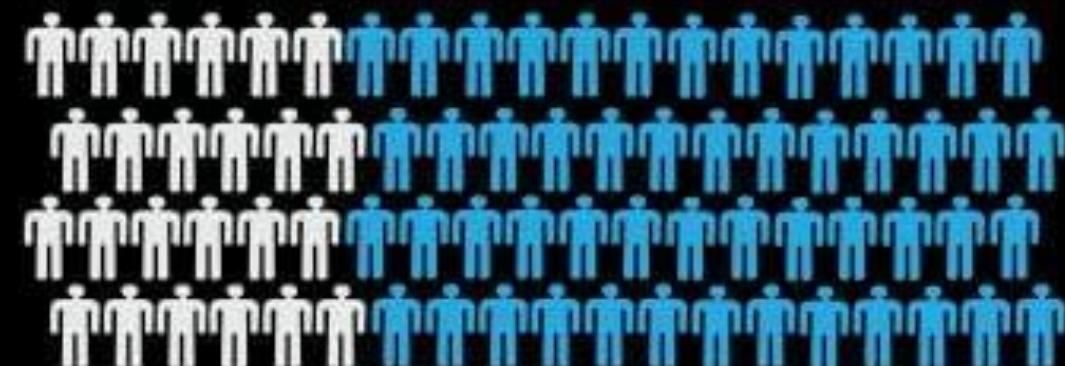
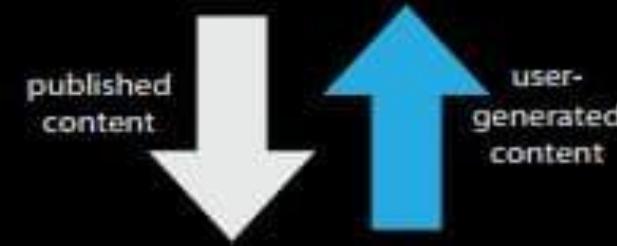
100,000,000 websites
(read-write Web)



1,000,000,000 users

Web 3.0

1,000,000,000 websites
(read-write Web)



2,500,000,000 users

Service-Oriented Computing (SOC)

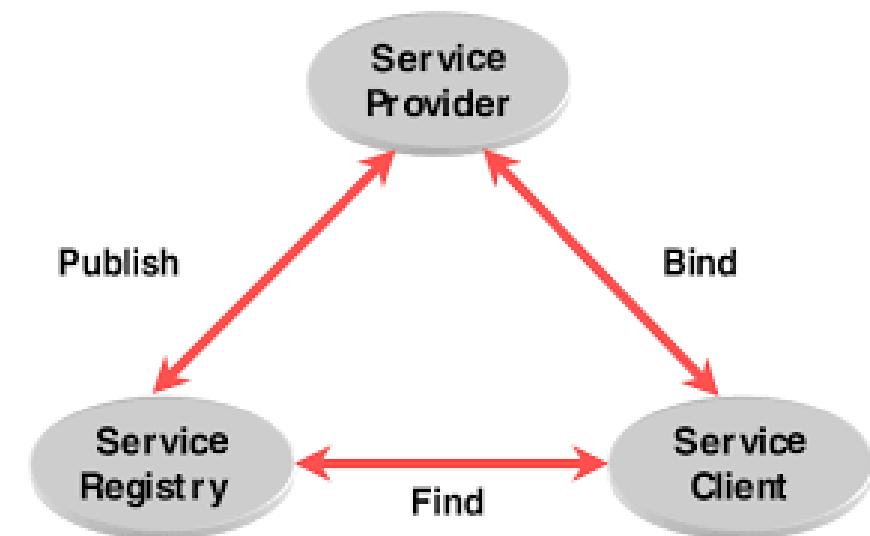
components into a network of services to create applications.

- Uses “services-oriented” programming to develop application by using network-available services.
- Web services are currently the most promising SOC-based technology. Uses internet-based standards:
 - Simple Object Access Protocol (SOAP)
 - Web Services Description Language (WSDL)
 - Business Process Execution Language for Web Services (BPEL4WS)

3

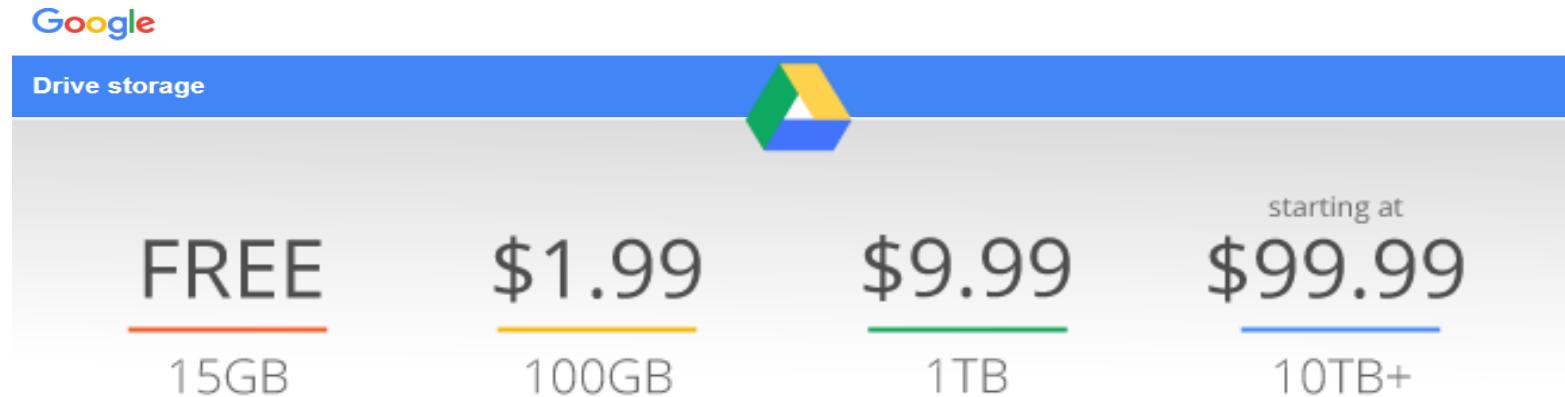
Service Oriented Computing

- Core Reference model for Cloud Computing System
- SOC Introduce Two main Concepts
 - Quality of Service (QoS)
 - Software as Service (SaaS)



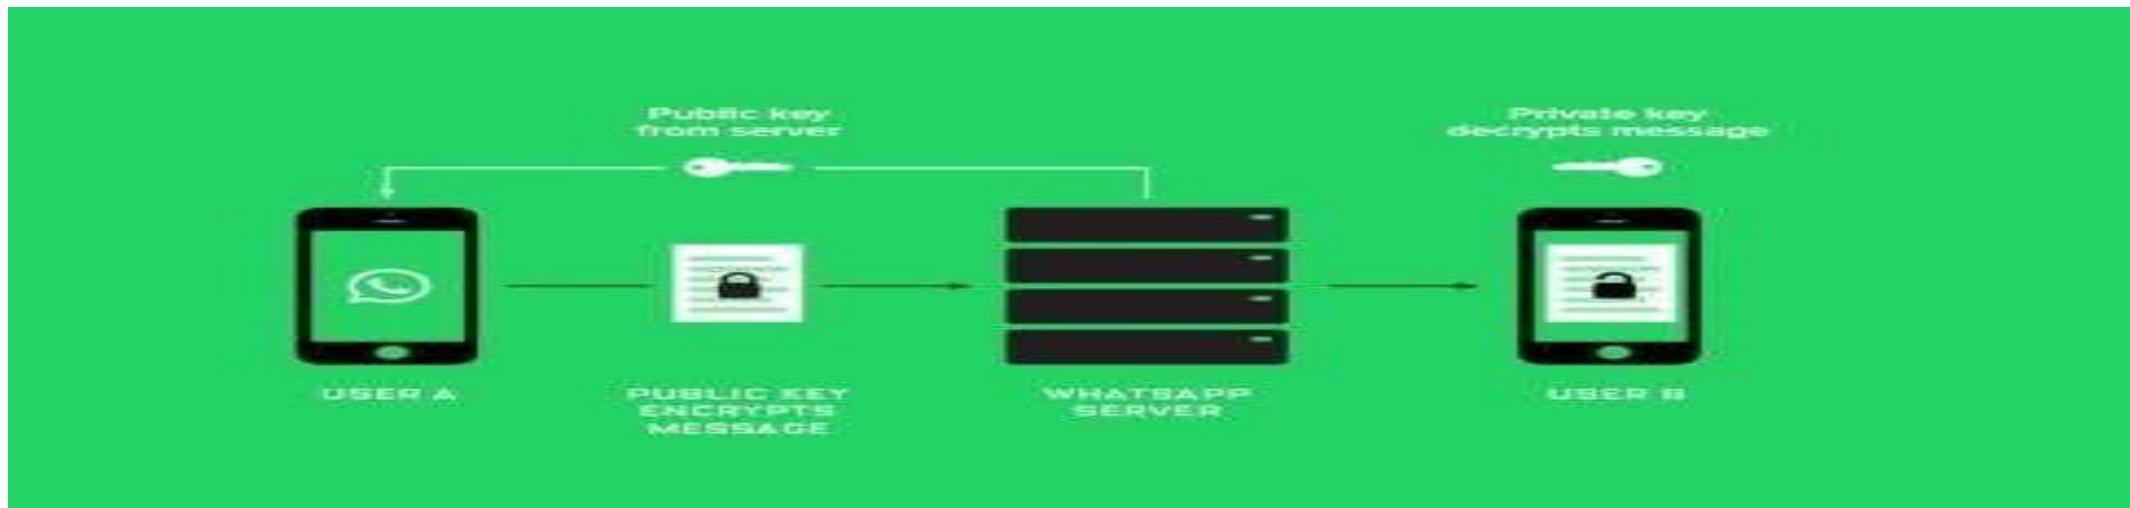
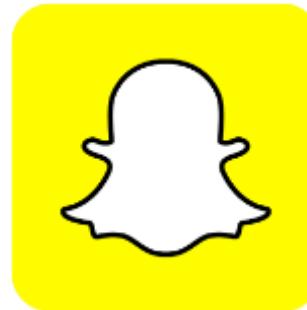
Utility Oriented Computing

- **The Computer Utility**, is a service provisioning model in which a service provider makes computing resources and infrastructure management available to the customer as needed, and charges them for specific usage rather than a flat rate.



Building Cloud Computing Environment

- Application Development



Enterprise Application



Gravitant

RIGHT SCALE®



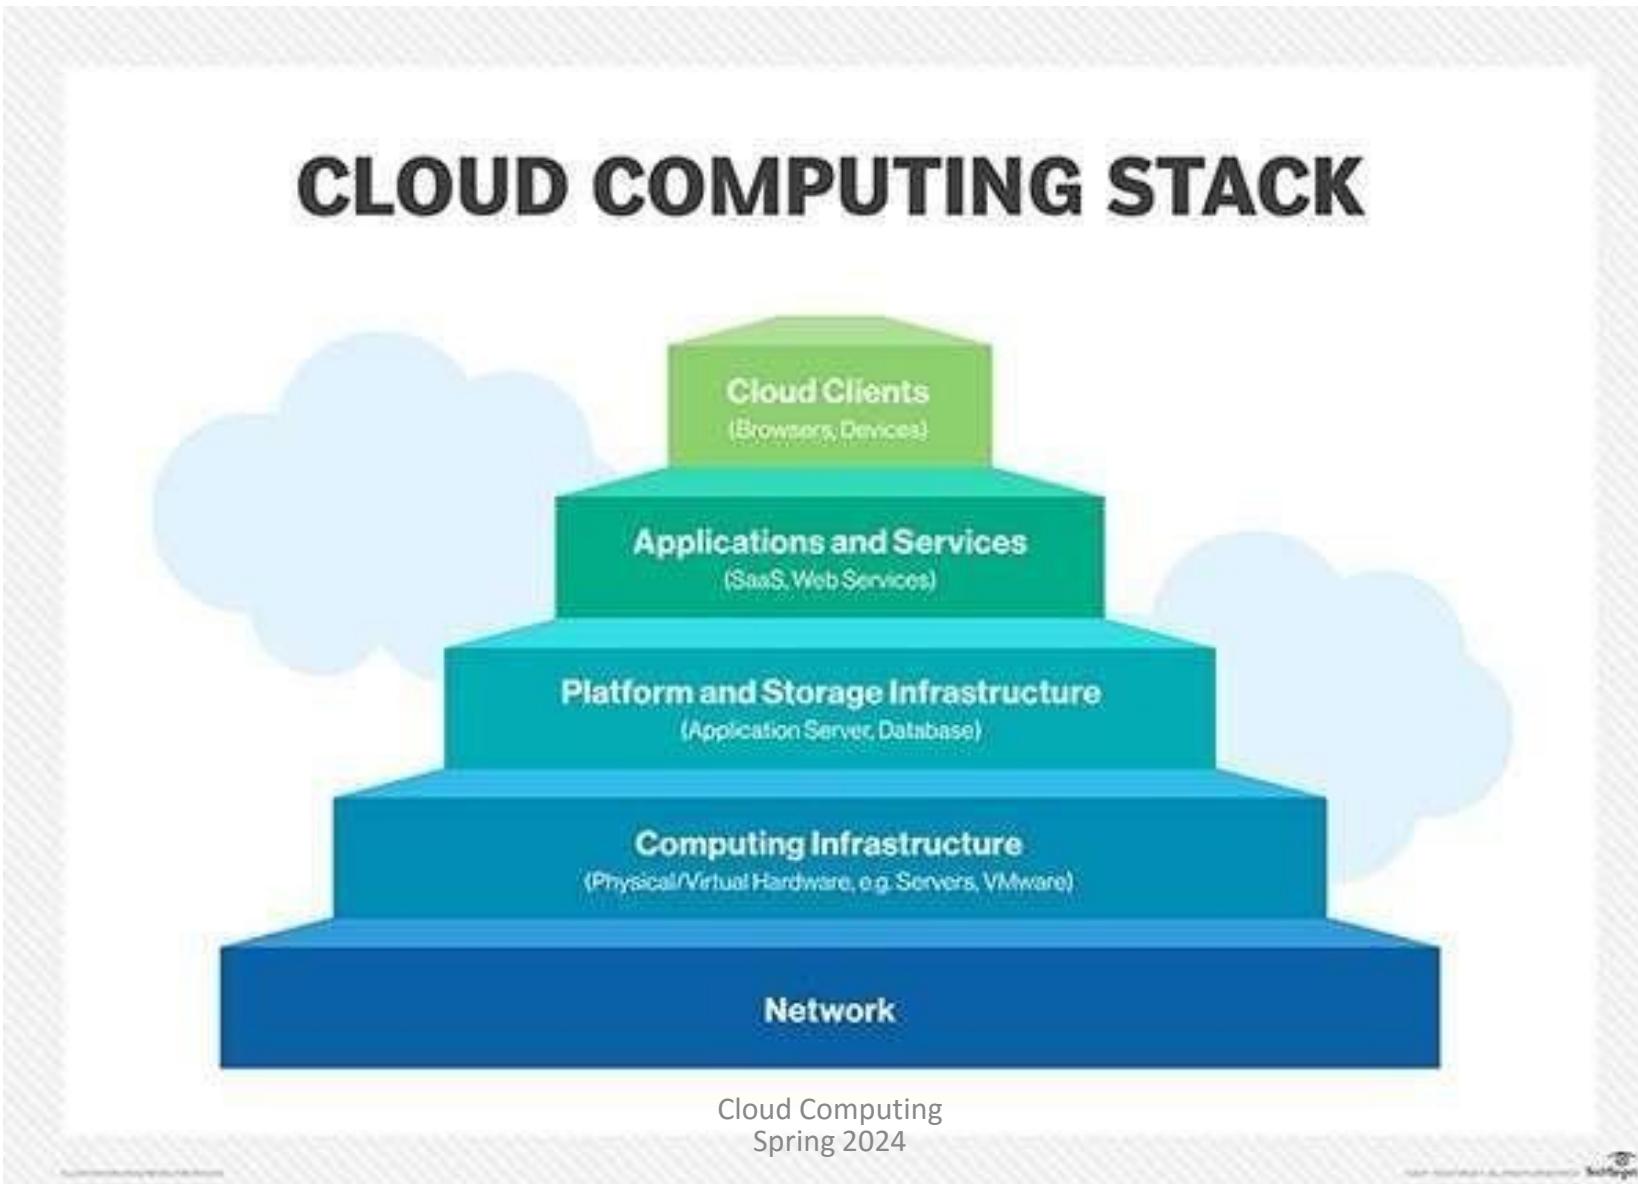
abiQUO®
ENTERPRISE CLOUD

enSTRATUS

servicemesh



Infrastructure and System Development



Computing Platform and Technologies

- AWS – Amazon Web Service



Free Cloud Computing Services - + Update :

aws.amazon.com/free/?nc2=h_ql_pr_ft&all-free-tier.sort-by=item.additionalFields.SortRank&all-free-tier.sort-order=asc&awsf.Free%20Tier%20... M

aws

Contact Us Support English My Account Sign In Create an AWS Account

Products Solutions Pricing Documentation Learn Partner Network AWS Marketplace Customer Enablement Events Explore More

AWS Free Tier Overview Free Tier Categories How to Create an Account Featured Offers for Business FAQs Terms and Conditions

AWS Free Tier

Gain free, hands-on experience with the AWS platform, products, and services

[Learn more about AWS Free Tier](#)

[Create a Free Account](#)

FEATURED

Startups may be eligible for AWS credits

AWS Activate provides eligible startups with a host of resources, including free AWS credits to spend on AWS services, and AWS Support.

[Sign up for Activate Today »](#)

Types Of Offers

Explore more than 100 products and start building on AWS using the Free Tier. Three different types of free offers are available depending on the product used. Click icon below to explore our offers.

Free trials

Short-term free trial offers start from the date you activate a particular service

12 months free

Enjoy these offers for 12-months following your initial sign-up date to AWS

Always free

These free tier offers do not expire and are available to all AWS customers

Activate Windows
Go to Settings to activate Windows.

Cloud Computing
Spring 2024

84

Google App Engine

- PaaS
- For Developers



App Engine

App Engine Application Platform X +

cloud.google.com/appengine

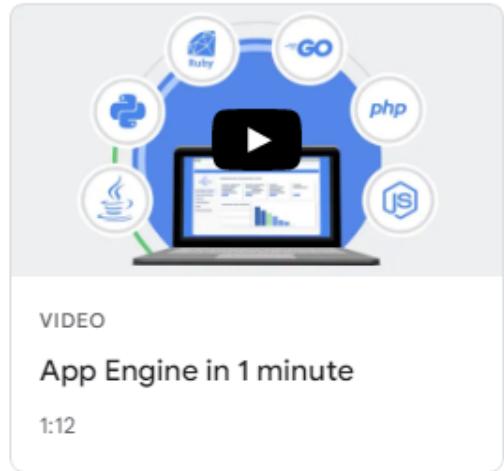
Google Cloud Overview Solutions Products Pricing Resources Docs Support Language Sign in

Contact Us Start free

App Engine

Build monolithic server-side rendered websites. App Engine supports popular development languages with a range of developer tools.

New customers get \$300 in free credits to spend on App Engine. All customers get 28 instances in standard environment free per day, not charged against your credits.



App Engine

Key features

Documentation

All features

Pricing

Take the next step

Try App Engine free

Contact sales

- ✓ Free up your developers with zero server management and zero configuration deployments
- ✓ Stay agile with support for popular development languages and a range of developer tools
- ✓ Explore more products in our [serverless](#) portfolio

Activate Windows
Go to Settings to activate Windows.

Debug

worker - 20160314t203733

default@1e41da2c92fa

Type a file name

cloud repository:/
▶ web
▼ worker
 README.md
 analyzer.js
 app.yaml
 package.json
 reddit.js
 server.js
 vision.js
 .gitignore
 LICENSE.md
 README.md

```

66  function analyze() {
67    let topicPromise = acquireTopic();
68    let redditPromise = reddit.getImageUrls();
69
70    Promise.all([topicPromise, redditPromise]).then((values) => {
71      let topic = values[0];
72      let urls = values[1];
73      let promises = [];
74      for (let url of urls) {
75        let p = vision.annotate(url).then(result) => {
76          return publishEvent(result, topic);
77        }).catch(err) => {
78          console.error('Error annotating event:' + util.inspect(err));
79        };
80        promises.push(p);
81      }
82      Promise.all(promises).then(() => {
83        // send a final event that lets the client know its done
84        publishEvent({
85          type: 'fin',
86          total: promises.length
87        }, topic).catch(err) => {
88          console.error('Error publishing fin event: ' + util.inspect(err));
89        });
90      });
91    }).catch(err) => {
92      console.error('Error processing images: ' + util.inspect(err));
93    });
94  }

```

Snapshots

vision.js:21

Variables

2016-03-14 (20:56:13)

| | |
|--------------|------------------------------|
| values | #<Array> |
| 0 | #<Topic> |
| name | projects/cloudcats-next/t... |
| pubsub | #<PubSub> |
| metadata | #<Object> |
| baseUrl | undefined |
| parent | #<PubSub> |
| id | projects/cloudcats-next/t... |
| createMethod | function (anonymous funct... |
| methods | #<Object> |
| interceptors | #<Array> |
| iam | #<IAM> |

Call Stack

| | |
|----------------------|----------------|
| (anonymous function) | vision.js:21 |
| annotate | vision.js:18 |
| (anonymous function) | analyzer.js:75 |

Microsoft Azure

- Paas
- <https://azure.microsoft.com>



Popular[AI + machine learning](#)[Analytics](#)[Compute](#)[Containers](#)[Databases](#)[DevOps](#)[Developer tools](#)[Hybrid + multicloud](#)[Identity](#)[Integration](#)[Internet of Things](#)[Management and governance](#)Popular

Explore some of the most popular Azure products

[Virtual Machines](#)

Provision Windows and Linux VMs in seconds

[Azure Virtual Desktop](#)

Enable a secure, remote desktop experience from anywhere

[Azure SQL](#)

Migrate, modernize, and innovate on the modern SQL family of cloud databases

[Azure Cosmos DB](#)

Build or modernize scalable, high-performance apps

[Azure Kubernetes Service \(AKS\)](#)

Deploy and scale containers on managed Kubernetes

[Azure Cognitive Services](#)

Add cognitive capabilities to apps with APIs and AI services

[App Service](#)

Quickly create powerful cloud apps for web and mobile

[Azure PlayFab](#)

Everything you need to build and operate a live game on one platform

[Azure Functions](#)

Execute event-driven serverless code functions with an end-to-end development experience

[Azure Quantum](#)

Jump in and explore a diverse selection of today's quantum hardware, software, and solutions

[Azure Arc](#)

Secure, develop, and operate infrastructure, apps, and Azure services anywhere

[Azure Operator Insights](#)

Remove data silos and deliver business insights from massive datasets

Digital Ocean



DigitalOcean Holdings, Inc. is an American multinational technology company and cloud service provider.

The company is headquartered in New York City, New York, USA, with 15 globally distributed data centers worldwide.

New! Premium CPU-Optimized Droplets are now available. Learn more →

We're hiring Blog Docs Get Support Sales



Products ^

Solutions ^

Marketplace

Community ^

Pricing

Log in

Sign up

FEATURED PRODUCTS

Droplets

Scalable virtual machines

Kubernetes

Managed Kubernetes clusters

Cloudways

Managed cloud hosting

App Platform

Get apps to market faster

Databases

COMPUTE

Droplets

Kubernetes

App Platform

Functions

CLOUD WEBSITE HOSTING

Cloudways

STORAGE

Spaces Object Storage

Volumes Block Storage

NETWORKING

Virtual Private Cloud (VPC)

Cloud Firewalls

Load Balancers

DNS

MANAGED DATABASES

MongoDB

MySQL

PostgreSQL

Redis™

DEVELOPER TOOLS

API

CLI

Support Plans

Monitoring

Uptime

Activate Windows

Hadoop

- **Hadoop** is an open source, Java-based programming framework that supports the processing and storage of extremely large data
- <https://cloud.google.com/hadoop>
- <http://hadoop.apache.org/>



The Apache™ Hadoop® project develops open-source software for reliable, scalable, distributed computing.

The Apache Hadoop software library is a framework that allows for the distributed processing of large data sets across clusters of computers using simple programming models. It is designed to scale up from single servers to thousands of machines, each offering local computation and storage. Rather than rely on hardware to deliver high-availability, the library itself is designed to detect and handle failures at the application layer, so delivering a highly-available service on top of a cluster of computers, each of which may be prone to failures.

[Learn more »](#)[Download »](#)[Getting started »](#)

Latest news

[Release 3.3.4 available](#)

2022 Aug 8

This is a release of Apache Hadoop 3.3 line.

Modules

The project includes these modules:

- **Hadoop Common:** The common utilities that support the other Hadoop modules.

Related projects

Other Hadoop-related projects at Apache include

- **Ambari™:** A web-based tool for provisioning, managing, and monitoring Apache Hadoop clusters

[Activate Windows](#)

Go to Settings to activate Windows.

Salesforce

- salesforce.com, inc. is an American cloud computing company headquartered in San Francisco, California.
- <https://www.salesforce.com/in/>

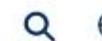




Products Industries Learning Support

Company Salesforce+

Contact Us
1800-420-7332



Login

Try for free

Customer 360 >

Sales

Service

Marketing

Commerce

Data Cloud

Tableau

Mulesoft

Slack

Platform

Net Zero

Salesforce can do for you.



Grow your small business

...and keep more customers using the business platform that grows with you.

LEARN MORE >



Questions?

No problem, we can help. Just start a chat with one of our sales agents, or call us at 1800-420-7332.

Activate Windows
LET'S CHAT NO THANKS >
Go to Settings to activate Windows.

What's new at Salesforce?

Cloud Computing
Spring 2024